



Stochastic Dual Coordinate Ascent Methods for Regularized Loss Minimization

Shai Shalev-Shwartz and Tong Zhang

School of CS and Engineering,
The Hebrew University of Jerusalem

"Optimization and Big Data",
Edinburgh 2013

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

Regularized Loss Minimization

$$\min_w P(w) := \left[\frac{1}{n} \sum_{i=1}^n \phi_i(w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \right].$$

Regularized Loss Minimization

$$\min_w P(w) := \left[\frac{1}{n} \sum_{i=1}^n \phi_i(w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \right].$$

Examples:

	$\phi_i(z)$	Lipschitz	smooth
SVM	$\max\{0, 1 - y_i z\}$	✓	✗
Logistic regression	$\log(1 + \exp(-y_i z))$	✓	✓
Abs-loss regression	$ z - y_i $	✓	✗
Square-loss regression	$(z - y_i)^2$	✗	✓

Dual Coordinate Ascent (DCA)

Primal problem:

$$\min_w P(w) := \left[\frac{1}{n} \sum_{i=1}^n \phi_i(w^\top x_i) + \frac{\lambda}{2} \|w\|^2 \right]$$

Dual problem:

$$\max_{\alpha \in \mathbb{R}^n} D(\alpha) := \left[\frac{1}{n} \sum_{i=1}^n -\phi_i^*(-\alpha_i) - \frac{\lambda}{2} \left\| \frac{1}{\lambda n} \sum_{i=1}^n \alpha_i x_i \right\|^2 \right]$$

- **DCA**: At each iteration, optimize $D(\alpha)$ w.r.t. a **single** coordinate, while the rest of the coordinates are kept in tact.
- **Stochastic** Dual Coordinate Ascent (**SDCA**): Choose the updated coordinate uniformly at random

SDCA vs. SGD — update rule

Stochastic Gradient Descent (SGD) update rule:

$$w^{(t+1)} = w^{(t)} - \eta_t \left(\lambda w^{(t)} + \phi'_i(w^{(t)\top} x_i) x_i \right)$$

SDCA update rule:

1. $\Delta_i = \operatorname{argmax}_{\Delta \in \mathbb{R}} D(\alpha^{(t)} + \Delta e_i)$
2. $\alpha_i^{(t+1)} = \alpha_i^{(t)} + \Delta_i$
3. $w^{(t+1)} = w^{(t)} + \frac{\Delta_i}{\lambda n} x_i$

- Often, the update rules are rather similar
- SDCA has several advantages:
 - Stopping criterion
 - No need to tune learning rate

SDCA vs. SGD — update rule — Example

SVM with the hinge loss: $\phi_i(w) = \max\{0, 1 - y_i w^\top x_i\}$

SGD update rule:

$$w^{(t+1)} = \left(1 - \frac{1}{t}\right) w^{(t)} - \frac{\mathbf{1}[y_i x_i^\top w^{(t)} < 1]}{\lambda t} x_i$$

SDCA update rule:

1. $\Delta_i = y_i \max \left(0, \min \left(1, \frac{1 - y_i x_i^\top w^{(t-1)}}{\|x_i\|_2^2 / (\lambda n)} + y_i \alpha_i^{(t-1)} \right) \right) - \alpha_i^{(t-1)}$
2. $\alpha_i^{(t+1)} = \alpha_i^{(t)} + \Delta_i$
3. $w^{(t+1)} = w^{(t)} + \frac{\Delta_i}{\lambda n} x_i$

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions


- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

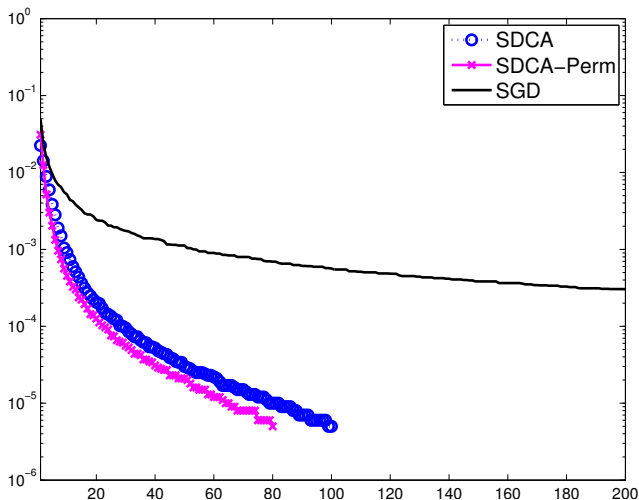
SDCA vs. SGD — experimental observations

- On CCAT dataset, $\lambda = 10^{-6}$, **smoothed loss**



SDCA vs. SGD — experimental observations

- On CCAT dataset, $\lambda = 10^{-6}$, **hinge-loss**



1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

SDCA vs. SGD — Current analysis is unsatisfactory

How many iterations are required to guarantee $P(w^{(t)}) \leq P(w^*) + \epsilon$?

- For SGD: $\tilde{O}\left(\frac{1}{\lambda\epsilon}\right)$
- For SDCA:
 - Hsieh et al. (ICML 2008), following Luo and Tseng (1992): $O\left(\frac{1}{\nu} \log(1/\epsilon)\right)$, but, ν can be arbitrarily small
 - S and Tewari (2009), Nesterov (2010):
 - $O(n/\epsilon)$ for general n -dimensional coordinate ascent
 - Can apply it to the dual problem
 - Resulting rate is slower than SGD
 - And, the analysis does not hold for logistic regression (it requires smooth dual)
 - Collins et al (2008): For smooth loss, similar bound to ours (for smooth loss) but for a more complicated algorithm (Exponentiated Gradient on dual)
 - Peter Richtarik and Martin Takac: Analysis of randomized coordinate descent with good rates (does not hold for logistic regression)
- Analysis is for **dual** sub-optimality

Dual vs. Primal sub-optimality

- Take data which is linearly separable using a vector w_0
- Set $\lambda = 2\epsilon/\|w_0\|^2$ and use the hinge-loss
- $P(w^*) \leq P(w_0) = \epsilon$
- $D(0) = 0 \Rightarrow D(\alpha^*) - D(0) = P(w^*) - D(0) \leq \epsilon$
- But, $w(0) = 0$ so $P(w(0)) - P(w^*) = 1 - P(w^*) \geq 1 - \epsilon$
- **Conclusion:** In the “interesting” regime, ϵ -sub-optimality on the dual can be meaningless w.r.t. the primal !

Very recent related work

- Lacoste-Julien, Jaggi, Schmidt, Pletscher (ICML 2013):
 - Study Frank-Wolfe algorithm for the dual of structured prediction problems.
 - Interestingly, boils down to SDCA for the case of binary hinge-loss
 - Same bound as our bound for the Lipschitz case
- Le Roux, Schmidt, Bach (NIPS 2012): A variant of SGD for smooth loss and finite sample.

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

Our results

- For $(1/\gamma)$ -smooth loss:

$$\tilde{O} \left(\left(n + \frac{1}{\lambda \gamma} \right) \log \frac{1}{\epsilon} \right)$$

- For L -Lipschitz loss:

$$\tilde{O} \left(n + \frac{L}{\lambda \epsilon} \right)$$


- For “almost smooth” loss functions (e.g. the hinge-loss):

$$\tilde{O} \left(n + \frac{1}{\lambda \epsilon^{1/(1+\nu)}} \right)$$

where $\nu > 0$ is a data dependent quantity

SDCA vs. DCA — Randomization is crucial


- On CCAT dataset, $\lambda = 10^{-4}$, smoothed hinge-loss



- In particular, the bound of Luo and Tseng holds for cyclic order, hence must be inferior to our bound


Smoothing the hinge-loss

$$\phi(x) = \begin{cases} 0 & x > 1 \\ 1 - x - \gamma/2 & x < 1 - \gamma \\ \frac{1}{2\gamma}(1 - x)^2 & \text{o.w.} \end{cases}$$




Smoothing the hinge-loss

- Mild effect on 0-1 error



Smoothing the hinge-loss

- Improves training time



- Duality gap as a function of runtime for different smoothing parameters

- Main lemma: for any t and $s \in [0, 1]$,

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{G^{(t)}}{2\lambda}$$

- $G^{(t)} = O(1)$ for Lipschitz losses
- With appropriate s , $G^{(t)} \leq 0$ for smooth losses

- Main lemma: for any t and $s \in [0, 1]$,

$$\mathbb{E}[D(\alpha^{(t)}) - D(\alpha^{(t-1)})] \geq \frac{s}{n} \mathbb{E}[P(w^{(t-1)}) - D(\alpha^{(t-1)})] - \left(\frac{s}{n}\right)^2 \frac{G^{(t)}}{2\lambda}$$

- **Bounding dual sub-optimality:**

Since $P(w^{(t-1)}) \geq D(\alpha^*)$, the above lemma yields a convergence rate for the dual sub-optimality

- **Bounding duality gap:** Summing the inequality for iterations $T_0 + 1, \dots, T$ and choosing a random $t \in \{T_0 + 1, \dots, T\}$ yields,

$$\mathbb{E} \left[(P(w^{(t-1)}) - D(\alpha^{(t-1)})) \right] \leq \frac{n}{s(T - T_0)} \mathbb{E}[D(\alpha^{(T)}) - D(\alpha^{(T_0)})] + \frac{sG}{2\lambda n}$$

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

Proximal version

Primal problem:

$$\min_w P(w) := \left[\frac{1}{n} \sum_{i=1}^n \phi_i(X_i^\top w) + \lambda g(w) \right]$$

Dual problem:

$$\max_{\alpha \in \mathbb{R}^{k \times n}} D(\alpha) := \left[\frac{1}{n} \sum_{i=1}^n -\phi_i^*(-\alpha_i) - \lambda g^* \left(\frac{1}{\lambda n} \sum_{i=1}^n X_i \alpha_i \right) \right]$$

Main Results

Assumptions:

- ϕ_i is L -Lipschitz or $1/\gamma$ -smooth w.r.t. $\|\cdot\|_P$
- g is strongly convex w.r.t. $\|\cdot\|_{P'}$
- $\|X_i\| \leq R$ where $\|X_i\| = \sup_{u \neq 0} \frac{\|X_i u\|_{D'}}{\|u\|_D}$

Results (bound on the number of iterations):

- For smooth losses: $\tilde{O} \left(\left(n + \frac{R^2}{\lambda\gamma} \right) \log(1/\epsilon) \right)$.
- For Lipschitz losses: $O \left(n + \frac{(RL)^2}{\lambda\epsilon} \right)$
- Approximate dual maximization suffices (can obtain closed form approximate solutions while still maintaining same guarantees on the convergence rate)

Structured Output Learning

- The optimization problem:

$$\min_w \left[\frac{\lambda}{2} \|w\|_2^2 + \frac{1}{n} \sum_{i=1}^n \left(\max_{y'} \delta(y', y_i) - w^\top \psi(x_i, y_i) + w^\top \psi(x_i, y') \right) \right] .$$

- Let $g(w) = \frac{1}{2} \|w\|_2^2$, let the j 'th column of X_i be $\psi(x_i, j)$ and let,

$$\phi_i(v) = \max_j (\delta(j, y_i) - v_{y_i} + v_j) .$$

- Then g is 1-strongly convex w.r.t. $\|\cdot\|_2$ and ϕ_i is 2-Lipschitz w.r.t. norm $\|\cdot\|_\infty$. Therefore,

$$\|X_i\| = \sup_{u \neq 0} \frac{\|X_i u\|_2}{\|u\|_1} = \sup_{u: \|u\|_1=1} \|X_i u\|_2 = \max_j \|\psi(x_i, j)\|_2 \leq R .$$

- The resulting algorithm is equivalent to Lacoste-Julien et al [2012].

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?

ℓ_1 regularization, instances of low ℓ_2 norm

- Solve:

$$\min_w \left[\frac{1}{n} \sum_{i=1}^n \phi_i(x_i^\top w) + \sigma \|w\|_1 \right],$$

- Assume: $R = \max_i \|x_i\|_2 = O(1)$
- Set $g(w) = \frac{1}{2} \|w\|_2^2 + \frac{\sigma}{\lambda} \|w\|_1$
- Then, $\nabla_i g^*(v) = \text{sign}(v_i) \left[|v_i| - \frac{\sigma}{\lambda} \right]_+$ corresponds to soft-thresholding
- Runtime (smooth case):

Ours	SGD	FISTA	Primal SCD
$nd + \frac{d}{\epsilon \sigma^2}$	$\frac{d}{\epsilon^2 \sigma^2}$	$\frac{nd}{\sqrt{\epsilon} \sigma}$	$\frac{dn}{\epsilon \sigma^2}$

ℓ_1 regularization, instances of low ℓ_∞ norm

- Now assume $\max_i \|x_i\|_\infty = R = O(1)$
- Use $g(w) = \frac{3 \log(d)}{2} \|w\|_q^2 + \frac{\sigma}{\lambda} \|w\|_1$, with $q = \frac{\log(d)}{\log(d)-1}$
- $\nabla_i g^*(v) = \begin{cases} \text{sign}(v_i) \left(a \left(|v_i| - \frac{\sigma}{\lambda} \right) \right)^{\frac{1}{q-1}} & \text{if } |v_i| > \frac{\sigma}{\lambda} \\ 0 & \text{otherwise} \end{cases}$
for some easy to calculate a
- Similar analysis as before

1 Stochastic Dual Coordinate Ascent (SDCA)

- Works great in practice
- Lack of adequate theoretical understanding
- New theoretical analysis + interesting implications

2 Proximal versions

- Structured Output Learning
- General regularizers (e.g. ℓ_1)

3 Acceleration and parallel implementation

- Interpolating between accelerated gradient and SDCA
- How to parallelize ?


Interpolating between accelerated gradient and SDCA

- SDCA rate $\tilde{O}\left(n + \frac{1}{\lambda}\right)$
- Nesterov's Accelerated (deterministic) Gradient Descent (AGD):
 $\tilde{O}\left(\frac{1}{\sqrt{\lambda}}\right)$
- ASDCA with mini-batches (of size m):

Algorithm	$\gamma\lambda n = \Theta(1)$	$\gamma\lambda n = \Theta(1/m)$	$\gamma\lambda n = \Theta(m)$
SDCA	n	nm	n
ASDCA	n/\sqrt{m}	n	n/m
AGD	\sqrt{n}	\sqrt{nm}	$\sqrt{n/m}$

Experimental demonstration

- On CCAT dataset, $\lambda = 1/n$, smoothed hinge-loss



How to parallelize ?

- n examples, d features, s computing nodes
- **Divide by examples:** each node gets n/s examples

How to parallelize ?

- n examples, d features, s computing nodes
- **Divide by examples:** each node gets n/s examples
- **Divide by features:** each node gets d/s features of all examples

How to parallelize ?

- n examples, d features, s computing nodes
- **Divide by examples**: each node gets n/s examples
- **Divide by features**: each node gets d/s features of all examples
- Runtime per iteration:

Algorithm	partition type	runtime	communication time
SDCA	features	d/s	$s \log^2(s)$
ASDCA	features	dm/s	$ms \log^2(s)$
ASDCA	examples	dm/s	$d \log(s)$
AGD	examples	dn/s	$d \log(s)$

How to parallelize ?

- n examples, d features, s computing nodes
- Assume: $\lambda n = \Theta(1)$
- Total runtime:

Algorithm	partition type	runtime	communication time
SDCA	features	dn/s	$ns \log^2(s)$
ASDCA	features	$dnm^{1/2}/s$	$nm^{1/2}s \log^2(s)$
ASDCA	examples	$dnm^{1/2}/s$	$ndm^{-1/2} \log(s)$
AGD	examples	$dn^{3/2}/s$	$n^{1/2}d \log(s)$

Summary

- SDCA works very well in practice
- So far, theoretical guarantees were unsatisfactory
- Our analysis shows that SDCA is an excellent choice in many scenarios
- Proximal versions
- Accelerated SDCA with mini-batches