

Object-Parallel Solution of Large-Scale Lasso Problems

Jonathan Eckstein

Rutgers University, NJ, USA

joint work with

György Mátyásfalvi (doctoral student)

Rutgers University, NJ, USA



Funded in part by US
National Science Foundation
Grant CCF-1115638

Why “Object-Parallel”?

- Current trend: growth in computer power is mainly through increased parallelism
 - Speed of executing a single thread is not improving much
- To solve huge problem instances, including from “big data”, we need scalably parallel algorithms and implementations
- The best ways to exploit parallelism are often dictated by problem instance data structure and are thus highly application-dependent

Suggests we need optimization environments in which

- The application guides the use of parallelism as flexibly as possible
- The underlying algorithms can adapt easily to very different data representations (across distributed memory)

An “Object-Parallel” Approach to Optimization

- An “inside-out” approach:
- User determines the most appropriate data layout for a given application
 - How variables are divided (and perhaps replicated) among processors
 - And similarly for constraints
 - User provides (presumably efficient) calculations of objectives, constraints, and their derivatives using these layouts
- The optimization code is an algorithmic “template” that is applied to the **user’s data representation**

What Kind of Optimization Algorithms Seem Suitable to Such a Context?

- Methods that use only first derivatives
 - Structure of Hessians may be too complex for users to contemplate
- Methods that do not directly solve linear systems:
 - Structure of **factors** of matrices like

$$\begin{bmatrix} H & J^T \\ J & 0 \end{bmatrix}$$

may be too much for users to think about, or may defeat simple parallelism within calculating gradients/Jacobians

- Restrictive, but not impossible:
 - These are largely the same restrictions that apply to most approaches to solving very large-scale problem instances

A Possible General NLP Approach (Digression): Relative-Error First-Order Approximate Augmented Lagrangians

- Quadratic augmented Lagrangian “outer loop”
- Subproblems created within this outer loop are either unconstrained or box-constrained
- Solve the subproblems using an efficient box-constrained first-order method
- Use an approximation criterion so that relatively little effort can be expended on each subproblem...
 - ...maybe just a handful of gradient / CG steps
- E & Silva (2013) suggest an algorithm of this form
 - Show global convergence in convex case
 - Encouraging empirical performance results (nonconvex)

Relative-Error Augmented Lagrangians: Background

$$\begin{array}{ll} \min & f(x) \\ \text{ST} & g_i(x) \leq 0 \quad i = 1, \dots, m_1 \\ & h_i(x) = 0 \quad i = 1, \dots, m_2 \\ & l \leq x \leq u \end{array}$$

or just

$$\begin{array}{ll} \min & f(x) \\ \text{ST} & g(x) \leq 0 \\ & h(x) = 0 \\ & l \leq x \leq u \end{array}$$

Basic algorithm:

$$x^k \in \arg \min_{x \in \mathcal{R}^n, l \leq x \leq u} \left\{ f(x) + \frac{1}{2c_k} \left(\sum_{i=1}^m \max \{0, p_i^{k-1} + c_k g_i(x)\}^2 + \|q^{k-1} + c_k h(x)\|^2 \right) \right\}$$

$$p_i^k = \max \{0, p_i^{k-1} + c_k g_i(x^k)\} \quad i = 1, \dots, m$$

$$q^k = q^{k-1} + c_k h(x^k)$$

- But don't want to invest too much effort in subproblem minimization when multiplier estimates are poor

- Let $L_c(x, p, q) = f(x) + \frac{1}{2c} \left(\sum_{i=1}^m \max \{0, p_i + c g_i(x)\}^2 + \|q + c h(x)\|^2 \right)$

Relative Error Criterion (E & Silva 2012)

Violation of KKT

$$y^k \in \partial_x L_{c_k}(x^k, p^{k-1}, q^{k-1})$$

$$\frac{2}{c_k} \left| \left\langle w^{k-1} - x^k, y^k \right\rangle \right| + \|y^k\|^2 \leq \sigma \left(\left\| \min \left\{ \frac{1}{c_k} p^{k-1}, -g(x^k) \right\} \right\|^2 + \|h(x^k)\|^2 \right)$$

Scalar $\sigma \in [0, 1)$

$$p_i^k = \max \left\{ 0, p^{k-1} + c_k g_i(x^k) \right\} \quad i = 1, \dots, m$$

Weird auxiliary sequence

$$q^k = q^{k-1} + c_k h(x^k)$$

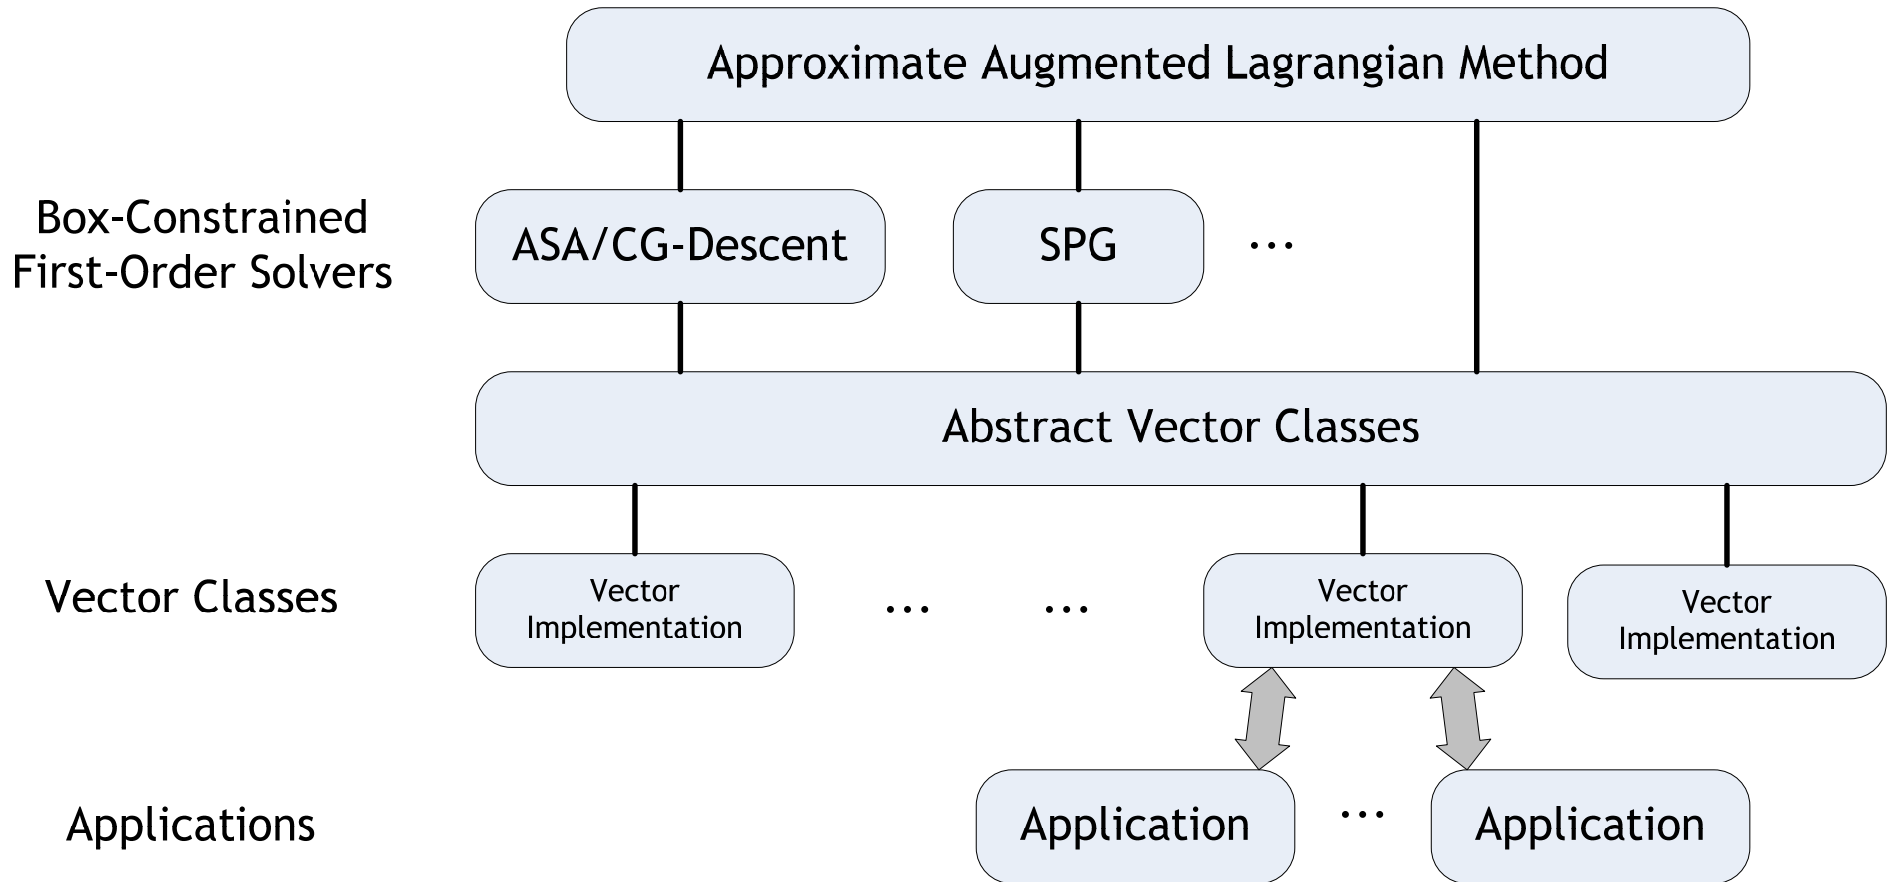
$$w^k = w^{k-1} - c_k y^k$$

The “FOAAL” Project

- The relative-error augmented Lagrangian approach seems like a reasonably good algorithm in serial — that is, likely to be competitive for at least some problem classes
- So, try to build this approach into object-parallel software:
- FOAAL: **F**irst-**O**rder **A**pproximate **A**ugmented **L**agrangian



FOAAL Software Architecture (C++)



FOAAL Software Architecture Rationale

- Vector classes encapsulate various vector representations
 - Serial representations of various kinds
 - ...including encapsulating various forms of BLAS
 - ...including parallel, such as Epetra (Heroux 2004, 2006)
- Abstract vector classes provide operator overloading so that the core optimization algorithms may be expressed with (almost) MATLAB-like simplicity
 - Unlike most high-performance optimization codes, in which the underlying algorithm is obscured by details
 - Without sacrificing much performance
 - Flexibly with respect underlying vector representation
 - Possibly in parallel (upper-level code running synchronously on all processors)

Operator Overloading

- When at least one of `a` and `b` has a user-defined type, C++ allows the definition of functions like `operator+` so that an expression like `a + b` implicitly calls `a.operator+(b)` etc.
- We use this feature to allow us to write simple MATLAB-like code for addition, subtraction, and scaling of vectors
- If implemented in the most straightforward way, could be extremely inefficient for large-scale (and parallel) vectors
 - Frequent creation and destruction of temporaries
 - Including parallel layout information
 - Unnecessary copying of memory
 - Excessive load and store operations

“Symbolic Temporaries” Technique

- Suppose `a` is a `double` and `w`, `x`, `y`, `z` are instances of our `VectorObject` class
- `VectorObject` encapsulates an `AbstractVector` pointer
- Consider an assignment statement like

$$w = x + a * y - z;$$

- At run time, this expression builds a relatively compact `LinearExpression` object expressing $(1, \&x), (a, \&y), (-1, \&z)$
- Overloading of the `=` operator invokes code to efficiently calculate the result and overwrite `w`
- Temporaries are still created, but they are compact and symbolic – their storage is independent of the vector length
- Overhead should be insignificant for large-scale applications (we have confirmed this)

“Symbolic Temporaries” Technique: More Details

$$w = x + a*y - z;$$

- A **LinearExpression** object is a list of pairs of the form (*scalar coefficient, pointer to vector*)
- Overload of **operator*(double&,VectorObject&)** produces a **LinearExpression** of the form (*a,&y*)
- Overload of **VectorObject.operator+(LinearExpression&)** appends (*1,&x*) to the **LinearExpression** list
- Overload of **LinearExpression.operator-(VectorObject&)** appends (*-1,&z*) to the **LinearExpression** list
- Overload of **VectorObject.operator=(LinearExpression&)** triggers actual calculation with minimum use of temporary memory (deferred evaluation technique)

The `LinearExpression` Class is Enough

First-order optimization methods typically require only a limited range of vector expressions

- Addition, scaling, (subtraction)
- Inner products (just produce a double)
- Projection on simple sets
- Function and gradient calculations:
 - We encapsulate these through a separate `Problem` class

Other operations occur, but are typically on scalars

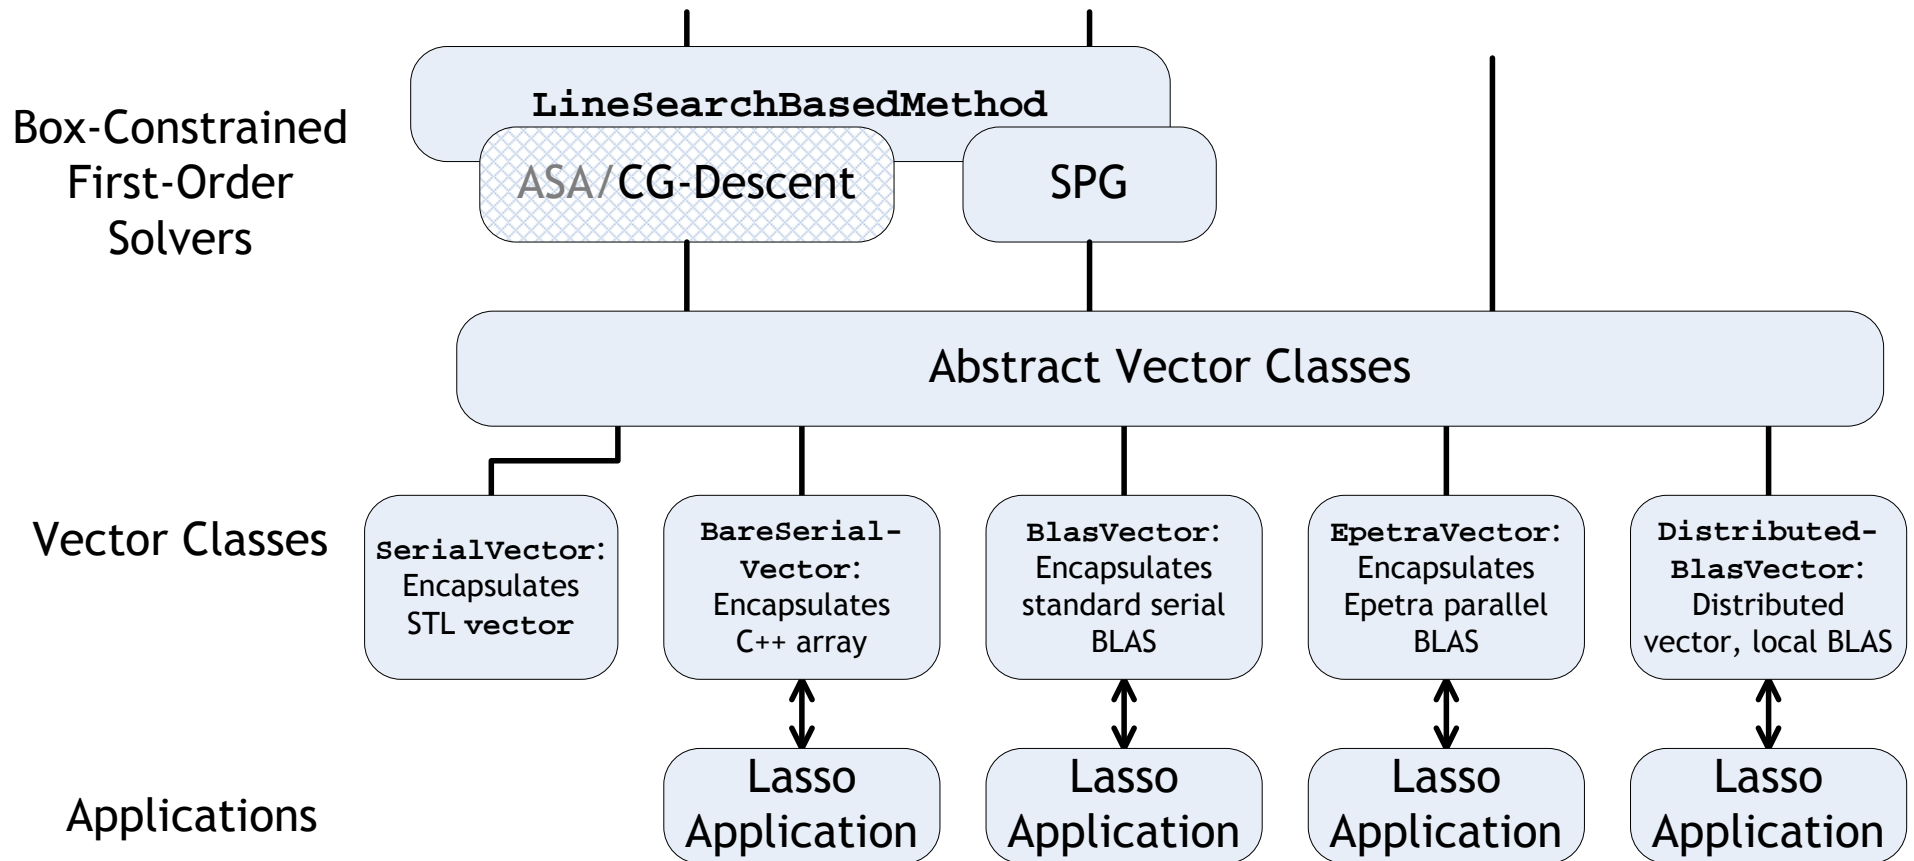
Support for Line-Search Methods

- We have also built some specialized support for line search
- Automatically cache function values and gradients to avoid both
 - Multiple evaluations at the same point, and...
 - Associated “clutter” in main algorithm code

What We Have Built So Far



- No top-level augmented Lagrangian yet
- Partial implementation of box solvers



- Lasso is a convenient (and potentially large-scale) test application for components already built

The Code for SPG Algorithm (with Nonmonotone Line Search)

```
VectorObject& min() {
    while(pGrad.norm_2() > tolerance && iter < maxIter) {
        iter++;
        xPlus = x + lambda*d;
        objValxPlus = Pr->objValGrad(xPlus, gk);
        objValMax = *(max_element(objValArray, objValArray + M));
        if(objValxPlus < objValMax + gamma*lambda*d.inner(g)) {
            descentSteps++;
            objValArray[descentSteps%M] = objValxPlus;
            xk = xPlus;
            b = inner(xk - x, gk - g);
            if(b <= 0) {
                stepSize = stepSizeMax;
            }
            else {
                a = norm_2_sq(xk-x);
                stepSize = smin(stepSizeMax, max(stepSizeMin, a/b));
            }
            x = xk;
            g = gk;
            xk = x - g;
            Pr->projectOnBounds(xk, pxk);
            pGrad = pxk - x;
            xk = x - stepSize*g;
            Pr->projectOnBounds(xk, pxk);
            d = pxk - x;
            lambda = 1;
        }
        else {
            objValx = Pr->objective(x);
            lambda = (-g.inner(d) * lambda*lambda )
                / (2*( objValxPlus - lambda*g.inner(d) - objValx ));
            if (lambda < 0.1 || lambda > 0.9*lambda)
                lambda = lambda/2;
        }
    }
}
```

Simple Test Application: Lasso

$$\boxed{\min \left\{ \frac{1}{2} \|Ax - b\|^2 + \nu \|x\|_1 \right\}} \quad (\text{unconstrained, nonsmooth})$$

Reformulated as

$$\boxed{\begin{array}{ll} \min & \frac{1}{2} \|A(x - y) - b\|^2 + \nu \cdot 1^T (x + y) \\ \text{ST} & x, y \geq 0 \end{array}}$$

- Becomes a box-constrained problem with a smooth objective
- In this form, does not test augmented Lagrangian level, but allows other software components to be tested
- Rest of talk will be about our experiences with (large-scale) Lasso

Overall Setup

- Fundamentally, using spectral projected gradient (SPG) on

$$\begin{array}{ll} \min & \frac{1}{2} \|A(x - y) - b\|^2 + \nu \cdot \mathbf{1}^\top (x + y) \\ \text{ST} & x, y \geq 0 \end{array}$$

- SPG implemented with our operator overloading technique
- Standard serial algorithm
- Parallelism through underlying linear algebra operations
 - Partly automatic in SPG / parallel vector classes
 - Partly application-specific (function and gradient evals)
- Main work per iteration is performing matrix multiplications of the form Ax and $A^\top u$:

- Objective gradient is $\nabla h(x, y) = \begin{bmatrix} A^\top (A(x - y) - b) + \nu \mathbf{1} \\ -A^\top (A(x - y) - b) + \nu \mathbf{1} \end{bmatrix}$

Data Distribution

- For now, assume A is large but has many more columns than rows ($n \gg m$) – fairly typical in Lasso
- Simple idea, sufficient for dense problems and some sparse problems: partition columns among processors, with matching partition of n -vectors:


$A_{(1)}$	\dots	$A_{(j)}$	\dots	$A_{(P)}$
$x_{(1)}$	\dots	$x_{(j)}$	\dots	$x_{(P)}$
$y_{(1)}$	\dots	$y_{(j)}$	\dots	$y_{(P)}$

- x and y are segments of a **DistributedBlasVector**
- Vectors of length m (like b) are replicated in every processor
 - Impact small because $n \gg m$

Matrix Multiplications

To multiply $z \in \mathbb{R}^n$ by A : $r = Az$


- Each processor performs a local multiplication $r_{(j)} = A_{(j)}z_{(j)}$
 - For Lasso, first compute $z_{(j)} = x_{(j)} + y_{(j)}$ (in parallel)
- Then do **Allreduce** operation (MPI primitive) to compute and broadcast $r = \sum_{j=1}^P r_{(j)}$
 - Requires $O(\log P)$ rounds of messages, time $O(m \log P)$



Transpose Matrix Multiplications

To multiply $p \in \mathbb{R}^m$ by A^\top : $g = A^\top p$

- Remember, m -vectors are replicated in each processor, so...
- Each processor performs a local multiplication $g_{(j)} = A_{(j)}^\top p$
- Result is g , properly distributed



Balancing Workload

- Key is for each of $A_{(1)}, \dots, A_{(P)}$ to have approximately the same number of nonzero elements
- Then the amount of work in computing $r_{(j)} = A_{(j)}z_{(j)}$, $g_{(j)} = A_{(j)}^T p$ will be roughly the same across processors, and the dominant portion of the computation will exhibit near-linear speedup
- Very easy for dense problems: just partition columns as equally as possible and relative imbalance will be $O(P/n)$
- One approach for sparse problems: heuristically partition the columns in order to balance the number of nonzero elements in each processor

Alternative Approach for Sparse Problems: Trust Epetra

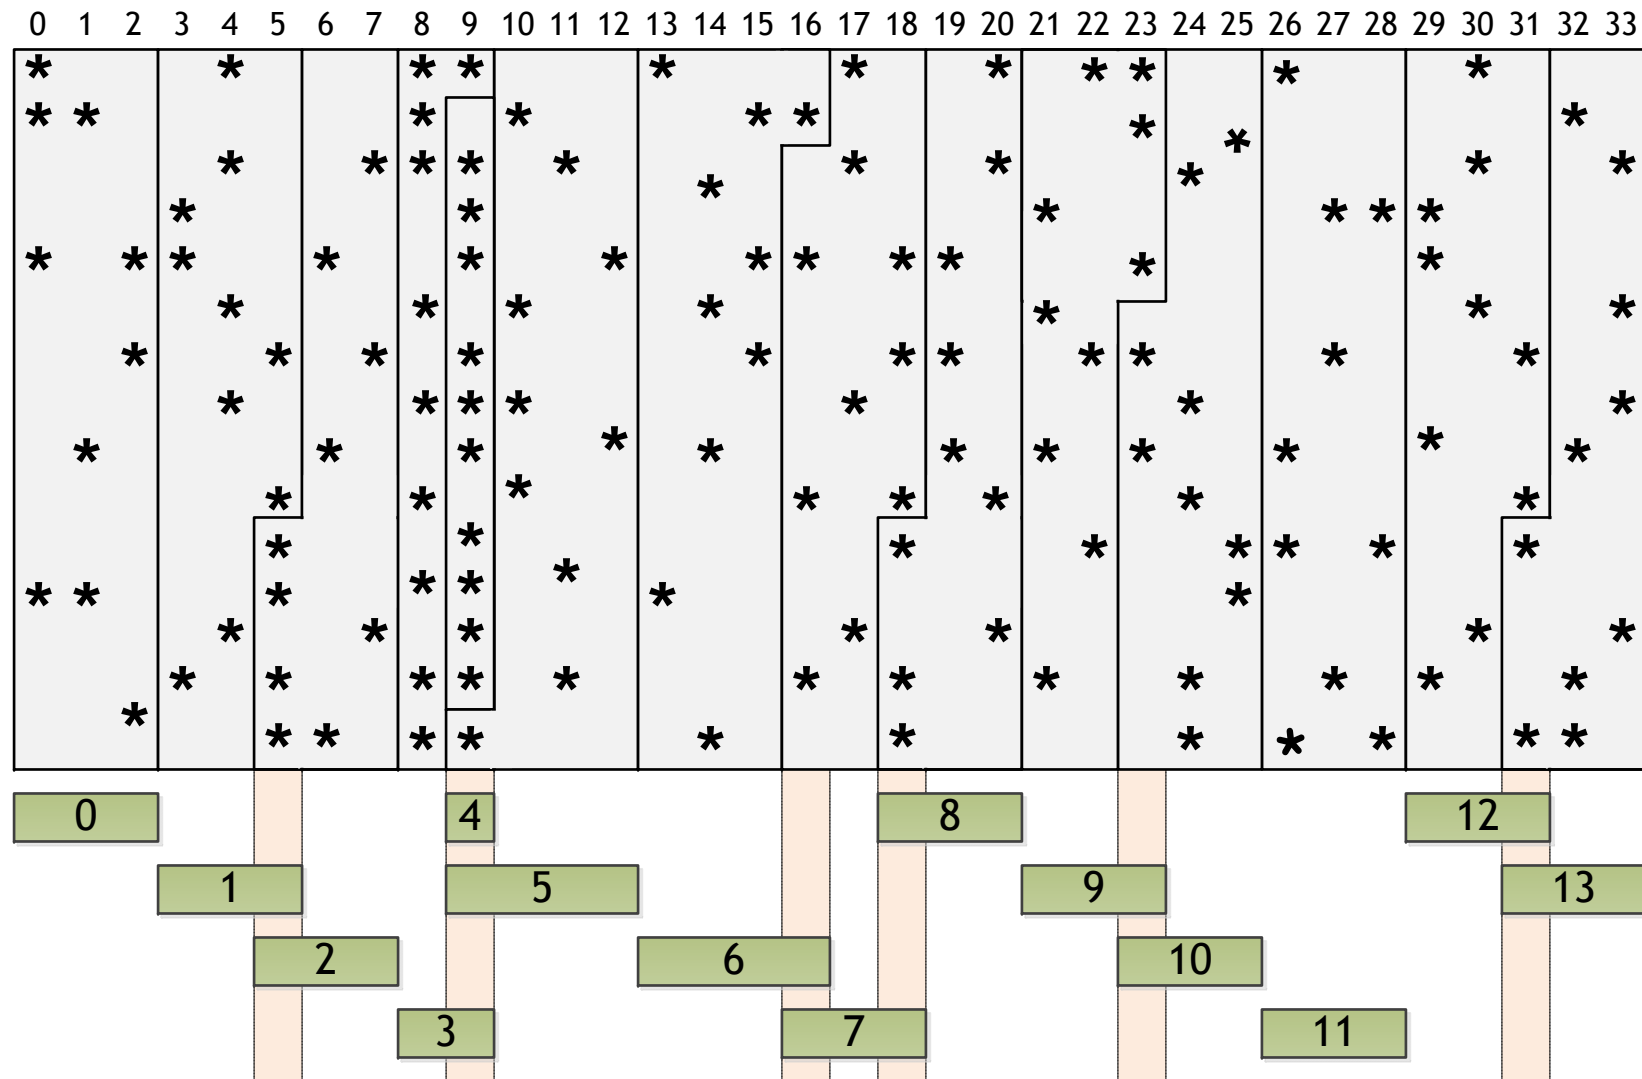
- Represent n -vectors using **EpetraVector**
- Partition matrices using Epetra's "Isorropia" sub-package
 - Can partition by rows and/or columns
 - Could be helpful for matrices that are closer to square
- Use Epetra parallel sparse matrix multiplication

Difficulties with Row/Column Partitioning

- We use matrices derived from real-world datasets
- Derived from UCI repository
- Many are sparse overall, but have a small fraction of relatively dense columns containing a significant fraction of the total nonzeros
- Rectangular partitions can work poorly such cases


Third Alternative for Sparse Problems

- Sort sparse matrix entries (*row*, *column*, *value*) by *column*
- Partition nonzeros evenly between processors, using this order




More on the Third Sparse Approach

- For n -vectors, each processor owns all the columns for which it has $(row, column, value)$ triplets
 - The first and/or last elements might overlap with adjacent processors (with replicated values)
- m -vectors are replicated on each processor as before
- Let $A_{(j)}$ denote the sparse matrix of locally stored triplets, filled out to rectangular shape by zeros
- The multiplication $r = Az$ works exactly the same as before



More on the Third Sparse Approach

- The multiplication $g = A^T p$ consists of the same local multiplication, followed by a scalar sum-reduction and broadcast within each “overlap zone”



Setting Things Up

- For implementation in MPI, we need to set up a “communicator” for each overlap zone
 - Once this is accomplished, implementation is easy
 - Reduction takes $O(\log V)$ steps, where $V \leq P$ is the maximum number of processors in an overlap zone
- Tricky to get communicators set up without introducing an $\Theta(P)$ operation, but can be done in $O(\log P)$

Digression: Parallel Scans

1.0	2.3	4.0	2.1	0.1	...
-----	-----	-----	-----	-----	-----



 →

1.0	3.3	7.3	9.4	9.5	...
-----	-----	-----	-----	-----	-----


- For any associative operator, this operation can be performed in $O(\log P)$ communication, $O((n/P)\log P)$ time
- Can design the operator to restart accumulation “from scratch” at certain points (implementing “segments”)

Digression: Details of Setting up Communicators

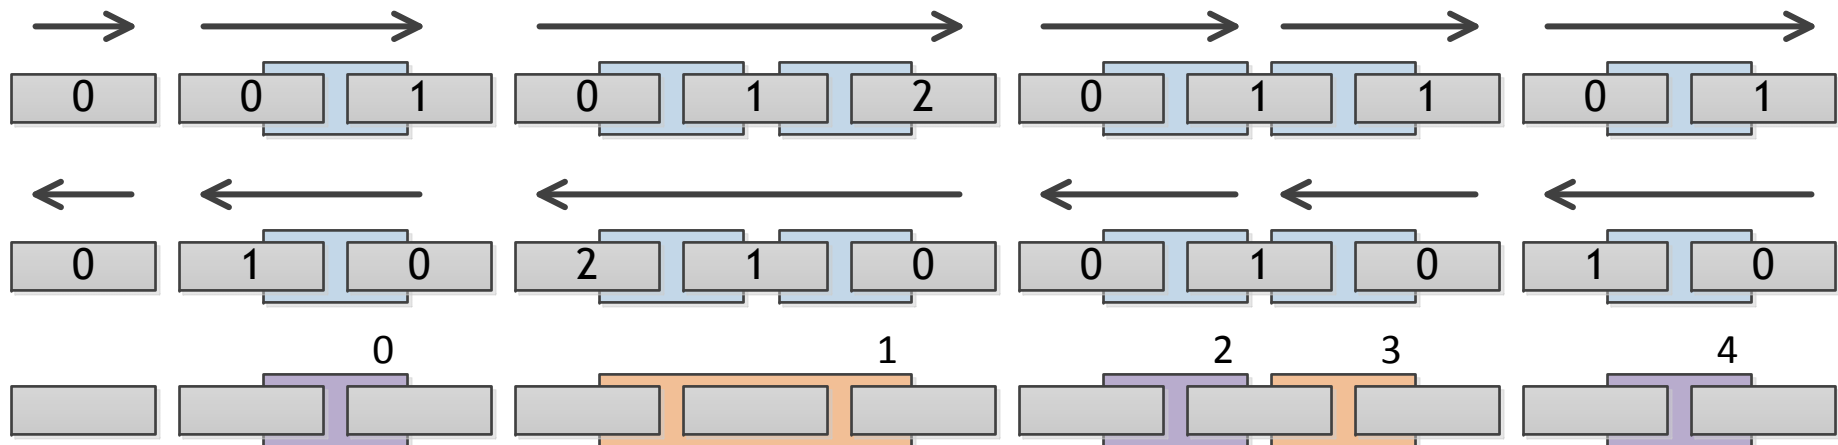
1. One message left, one message right to identify neighbors



2. Sum-scan to assign a unique number to each overlap group



3. Scan forward and back with a customized operator so each processor knows full extent of left and right group




Continued Digression: Detail of Setting Up Communicators

4. Call `MPI_Comm_create` to make all even-numbered groups



5. Call `MPI_Comm_create` to make all odd-numbered groups



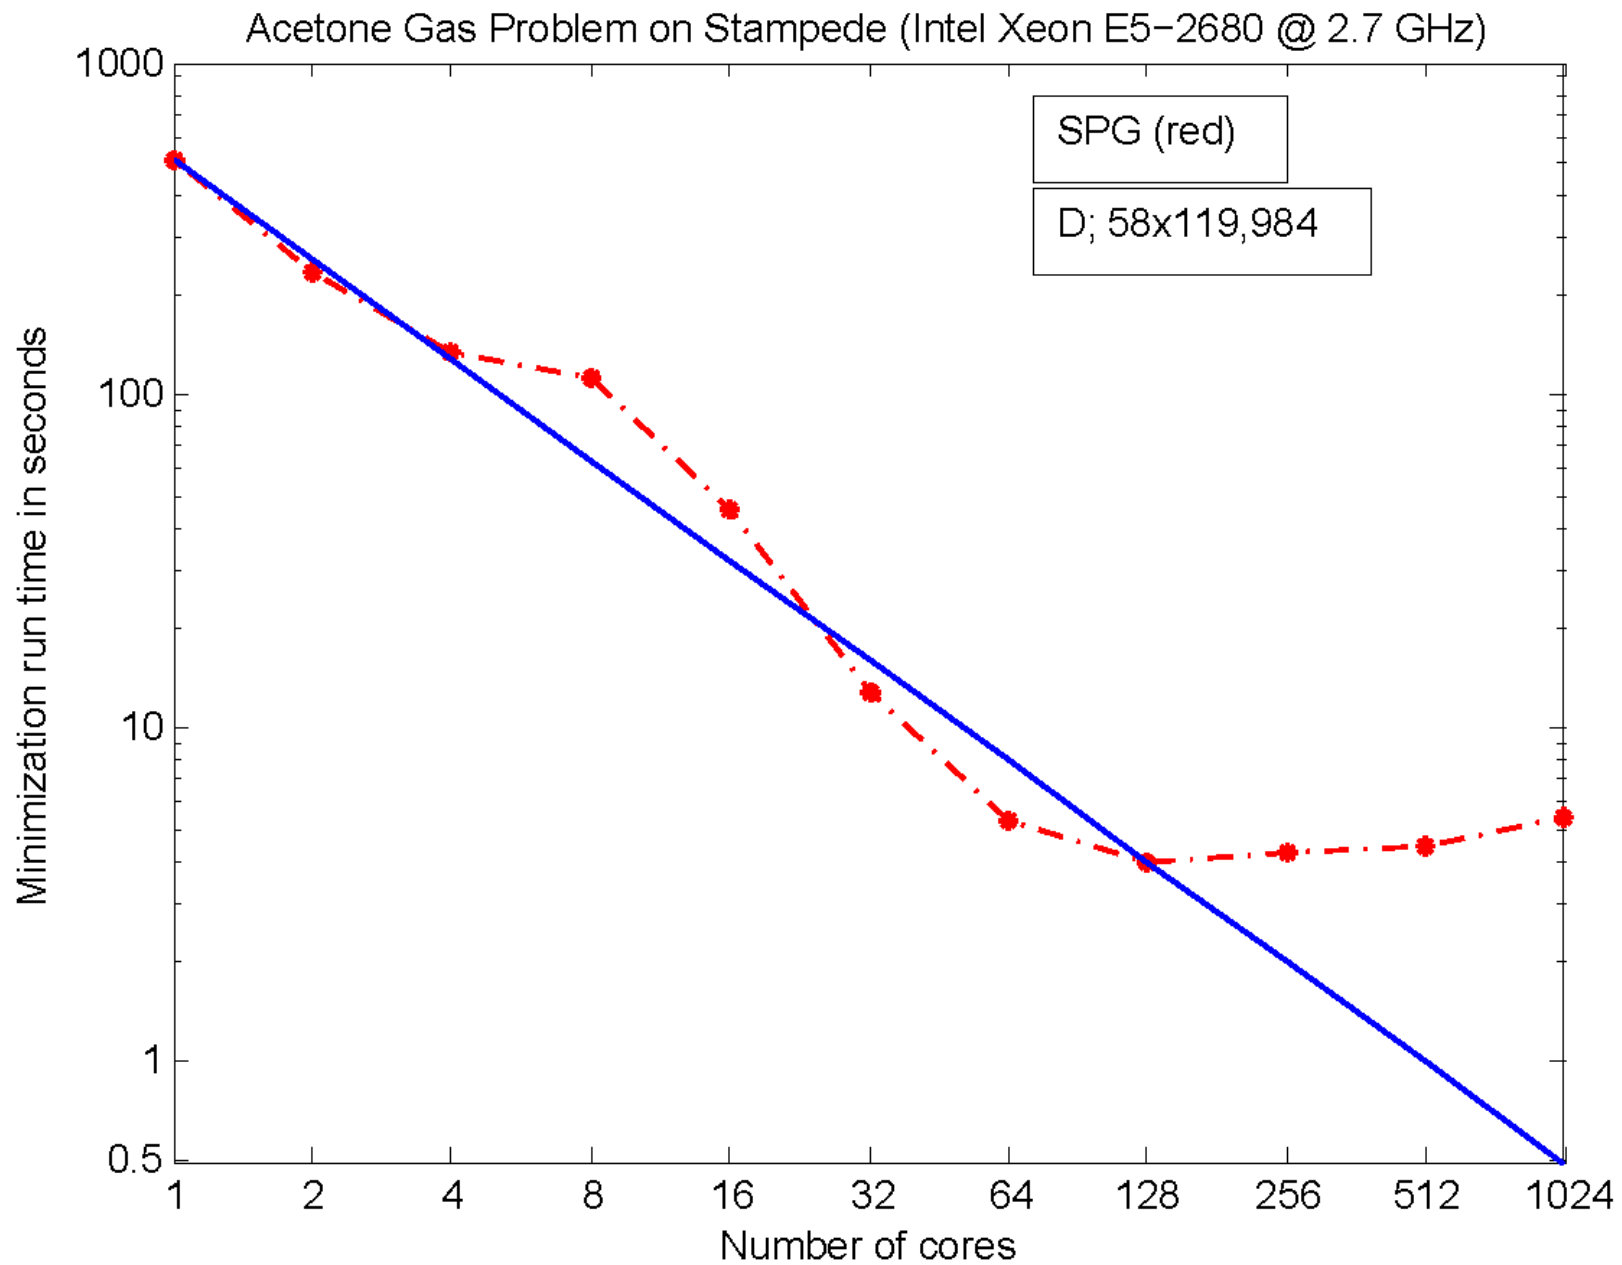
- Each call subdivides an existing communicator
- Each processor in at most one communicator
- All processors call `MPI_Comm_create` synchronously
- Those with nothing to do create and destroy a trivial communicator

This procedure avoids using tempting routines like `MPI_Comm_split`, which perform gather / scatter operations that can introduce $O(P)$ communication, hurting scalability

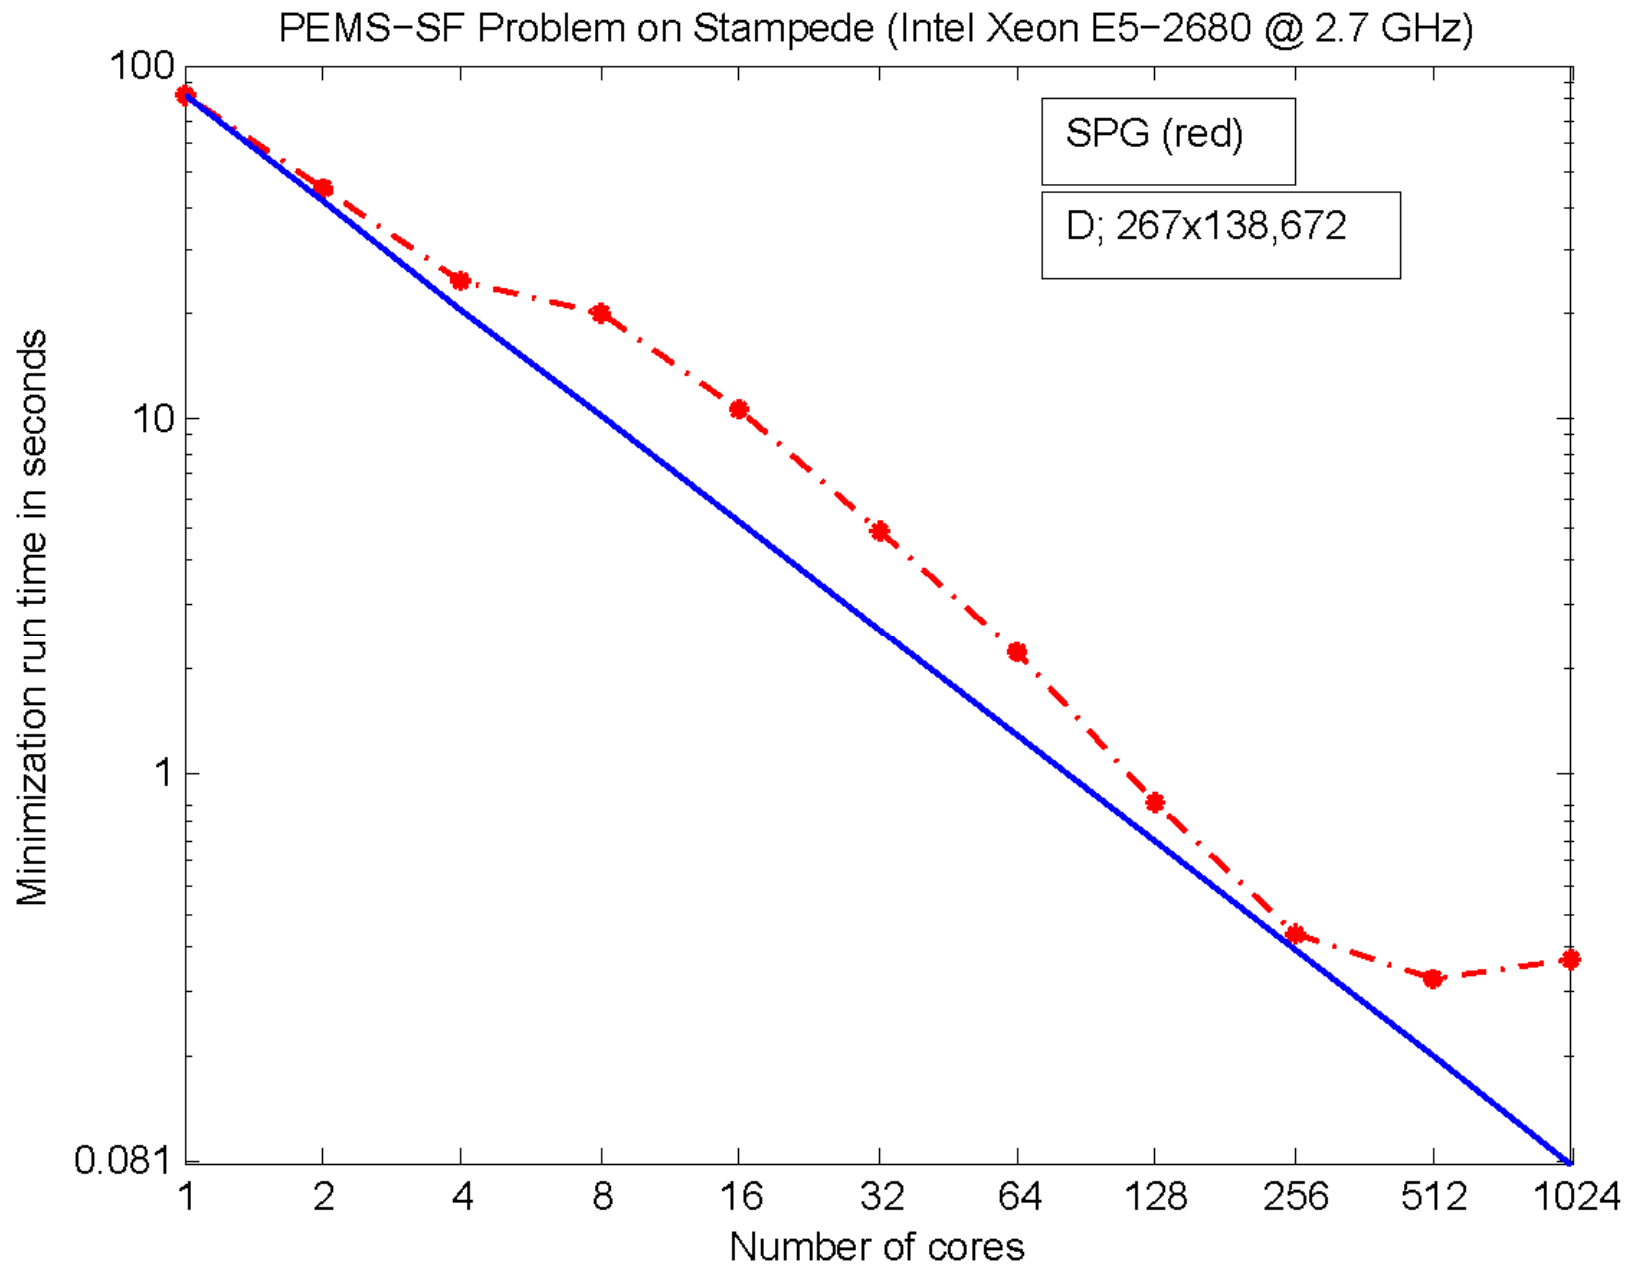
Computational Experiments

- SPG is a numerically sensitive algorithm
 - Minor changes (even version of BLAS) can cause some variation (not gigantic) in the number of iterations
- Convergence condition: $\left\| \left[(x, y) - \nabla h(x, y) \right]_+ - (x, y) \right\|_\infty \leq 10^{-6}$
- Some data adapted from UCI machine learning Repository
- Other data randomly generated
- Computer system: TACC Stampede
 - Xeon E5 cores, 2.7 GHz clock
 - Xeon Phi accelerators, but not used by our code
 - Infiniband fat-tree interconnect, MPI

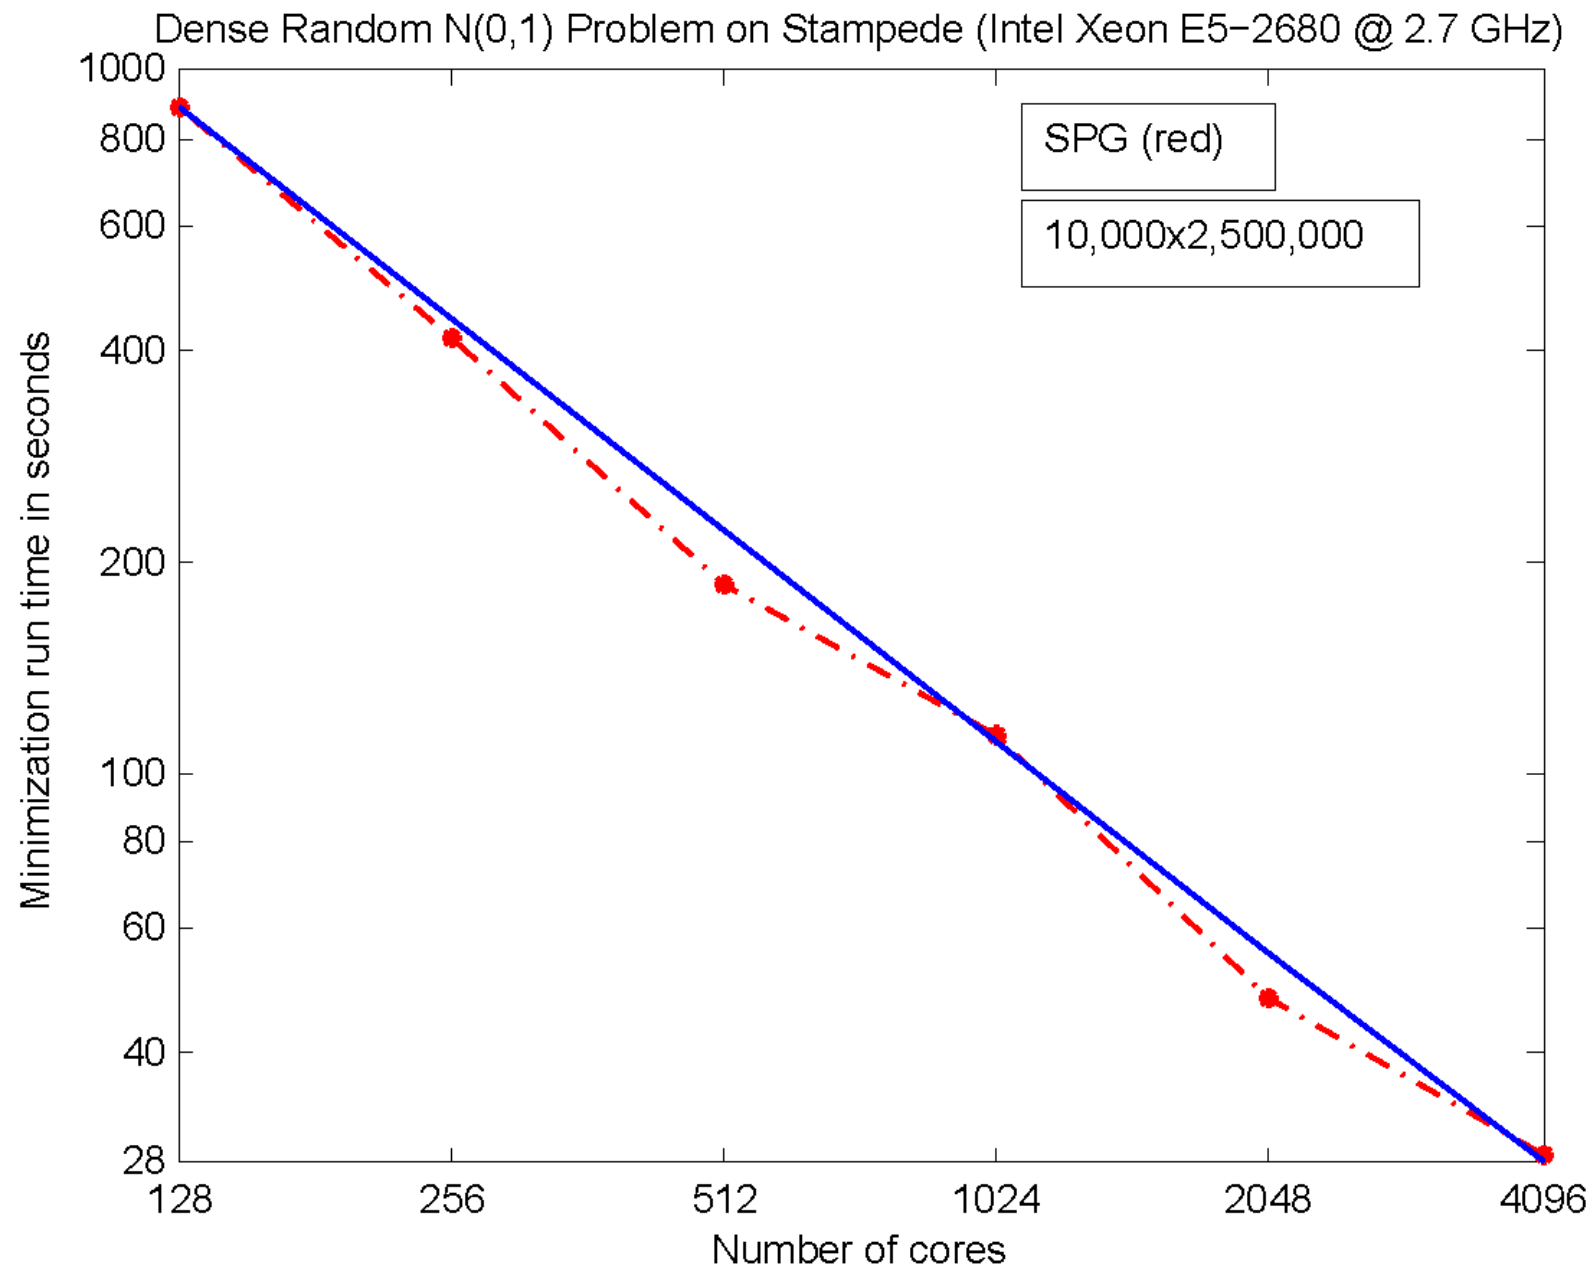
Dense: Detecting Acetone in Gas Samples (from UCI Repository)



More Dense: Highway Lane Occupancy (UCI Repository)



Dense: Randomly Generated Data



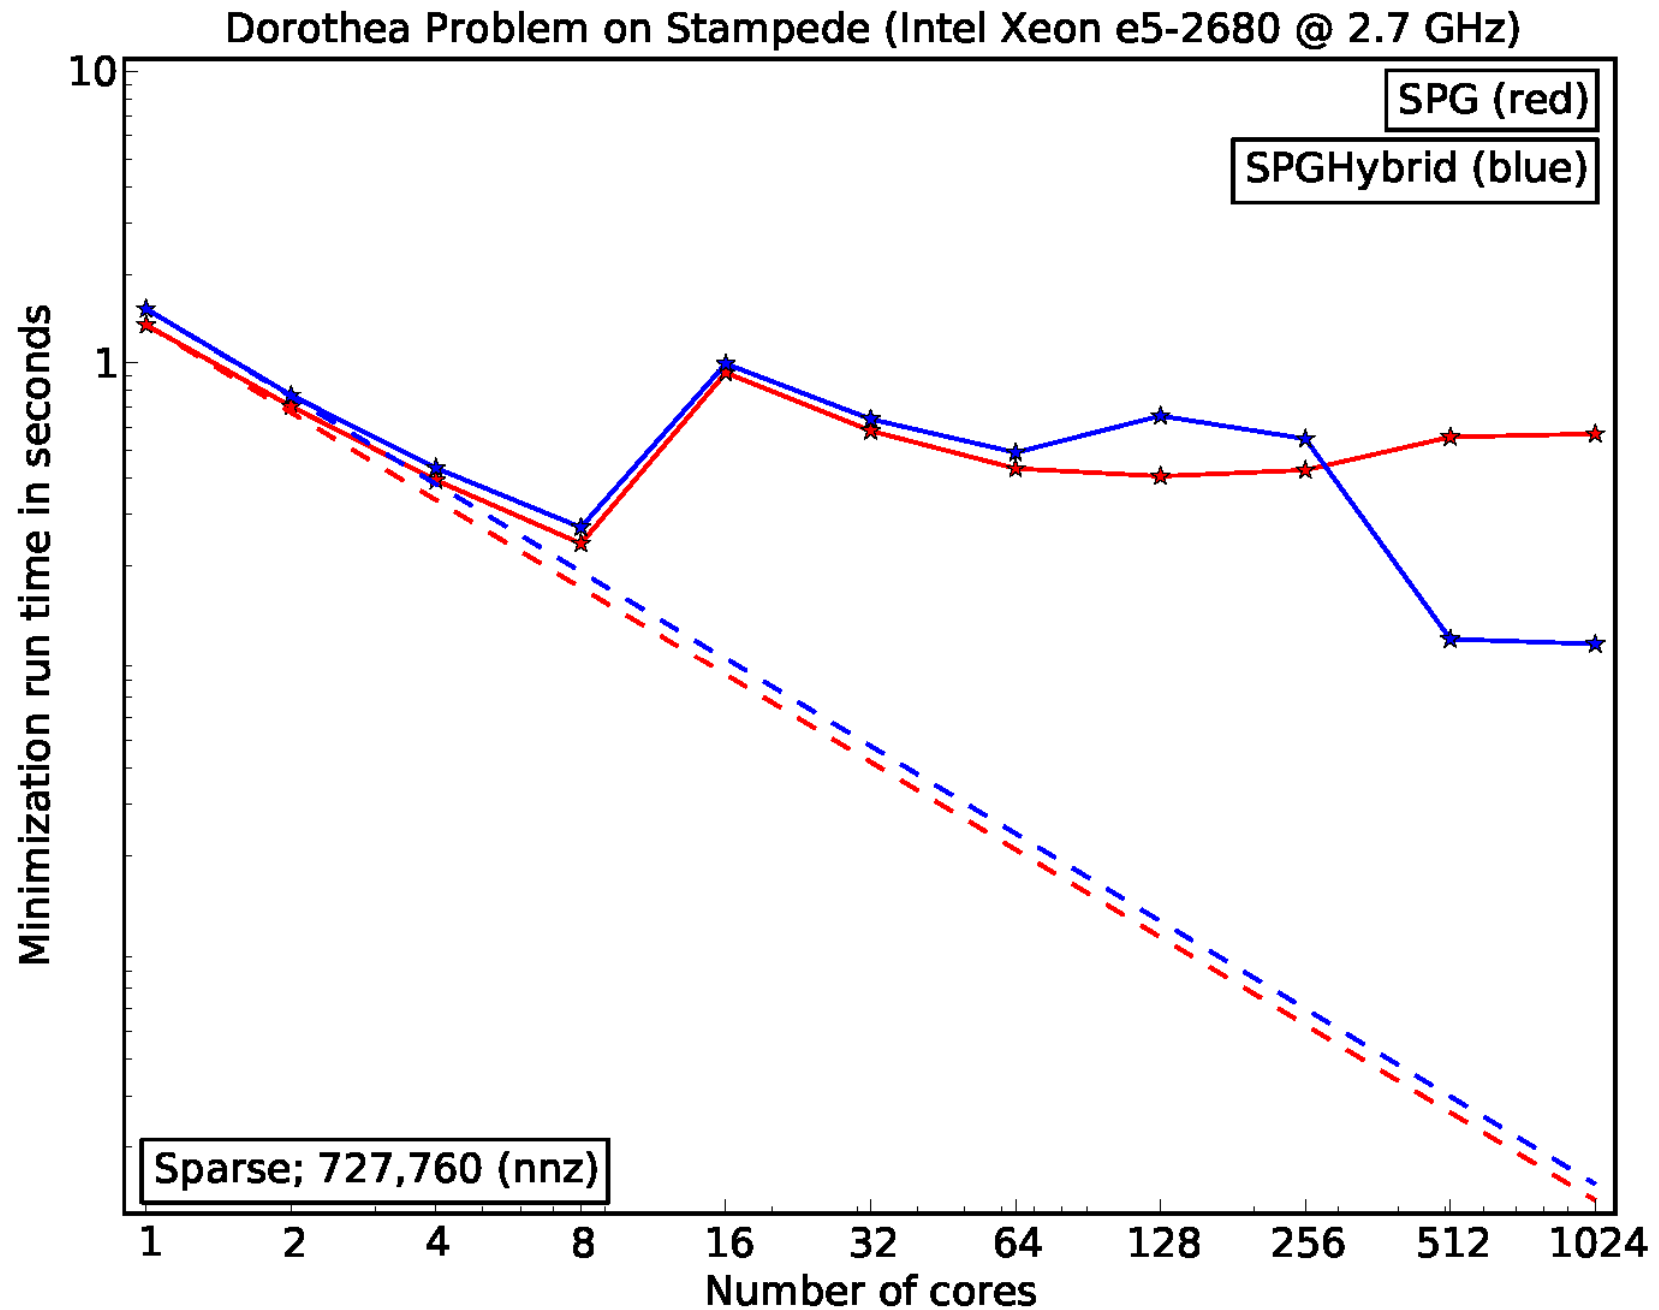
Sparse Problems

- “SPG” : our method with balanced nonzero partition
- “SPGHybrid” : our method with heuristic column partition
- “SPGEpetra” : our method with Epetra partitioning of matrix and vectors
- “ACDC” : Richtárik / Takáč random parallel coordinate descent
 - With small setting of proximal parameter - far below threshold required by (probabilistic) convergence theory
 - Does not have its own termination test; we terminate when it is close to the optimal value computed by SPG

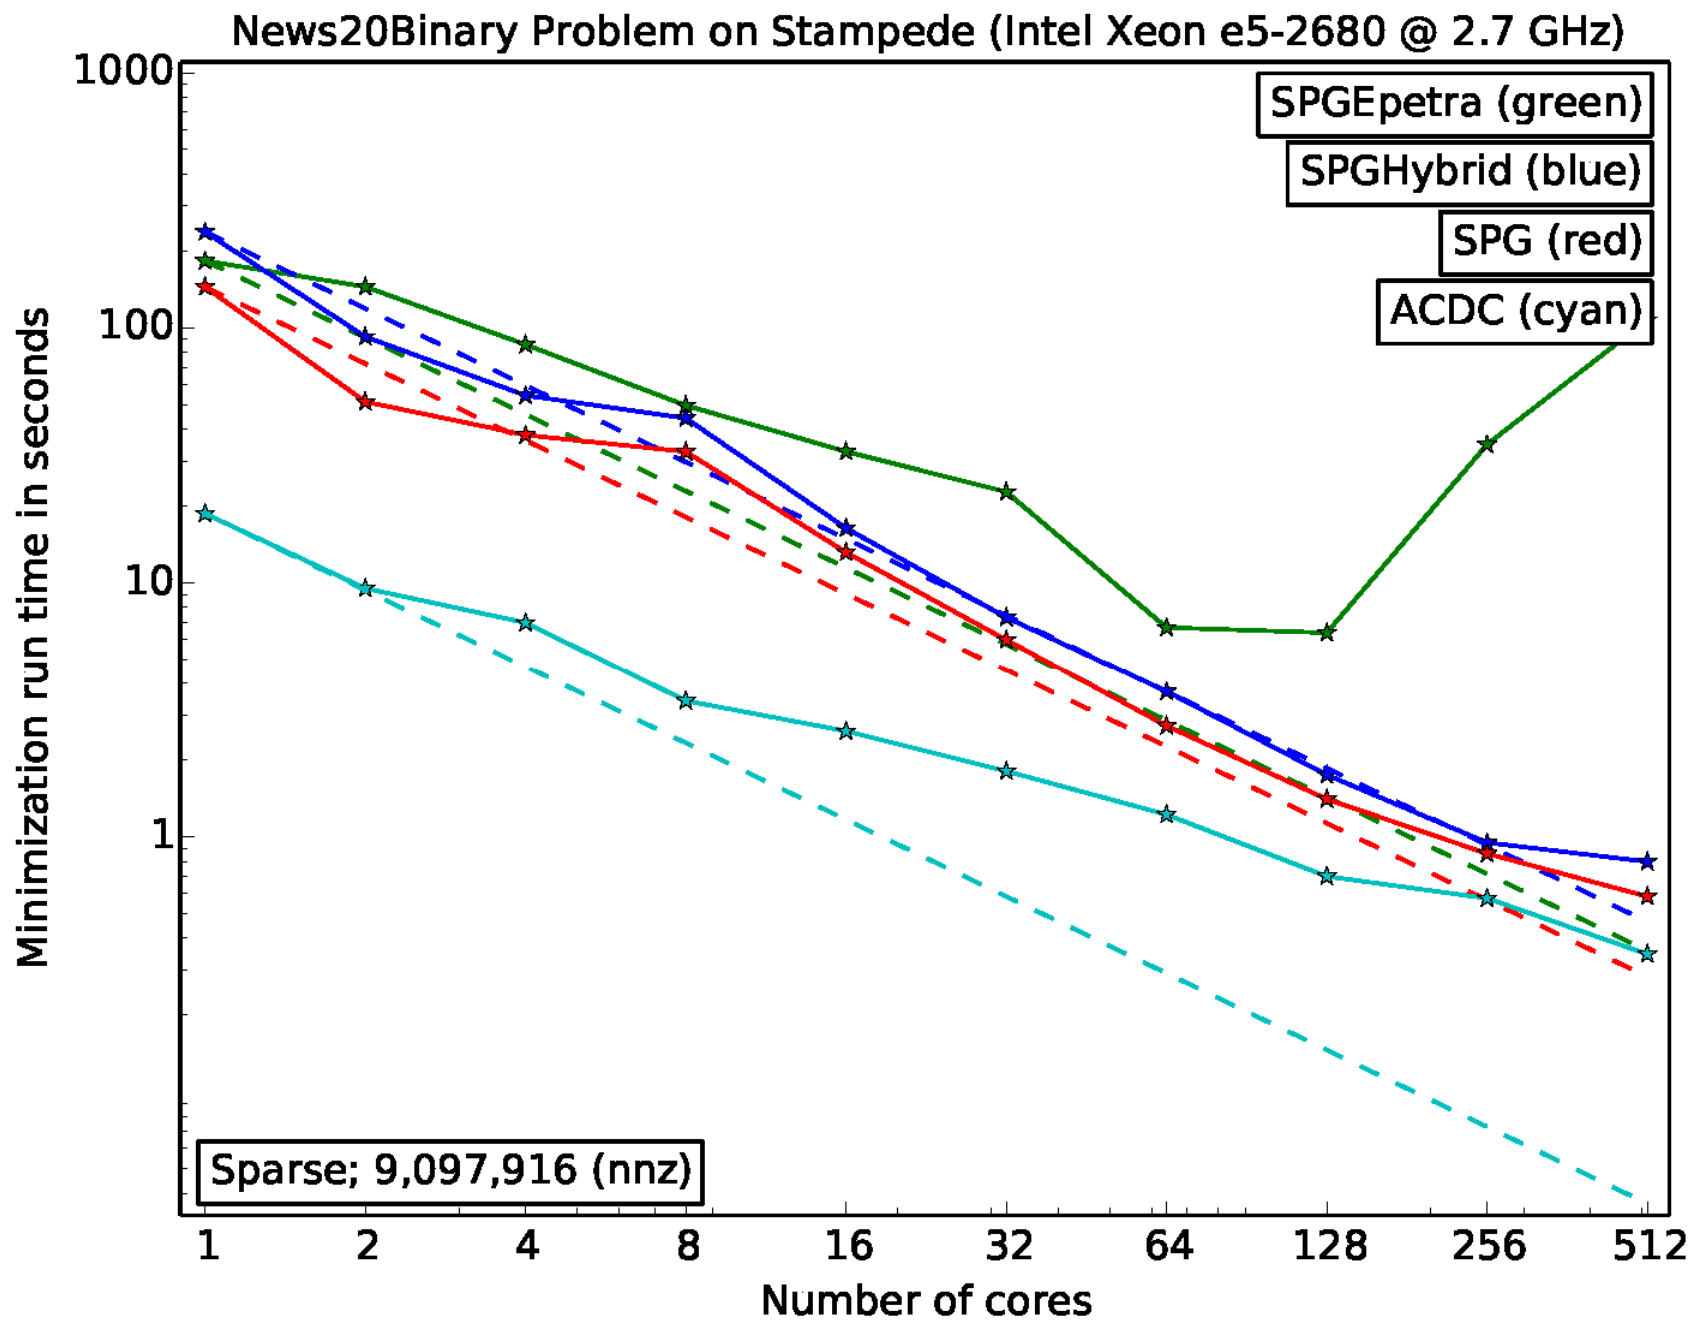
Current datasets: (working on getting more)

- Drug discovery: $1,950 \times 100,000$, 728K nonzeros
- News articles: $19,996 \times 1,355,191$, 9.1M nonzeros
- Randomly generated: $10,000 \times 2,500,000$, 250M nonzeros

Sparse: Drug Discovery (Small)




Sparse: News Articles



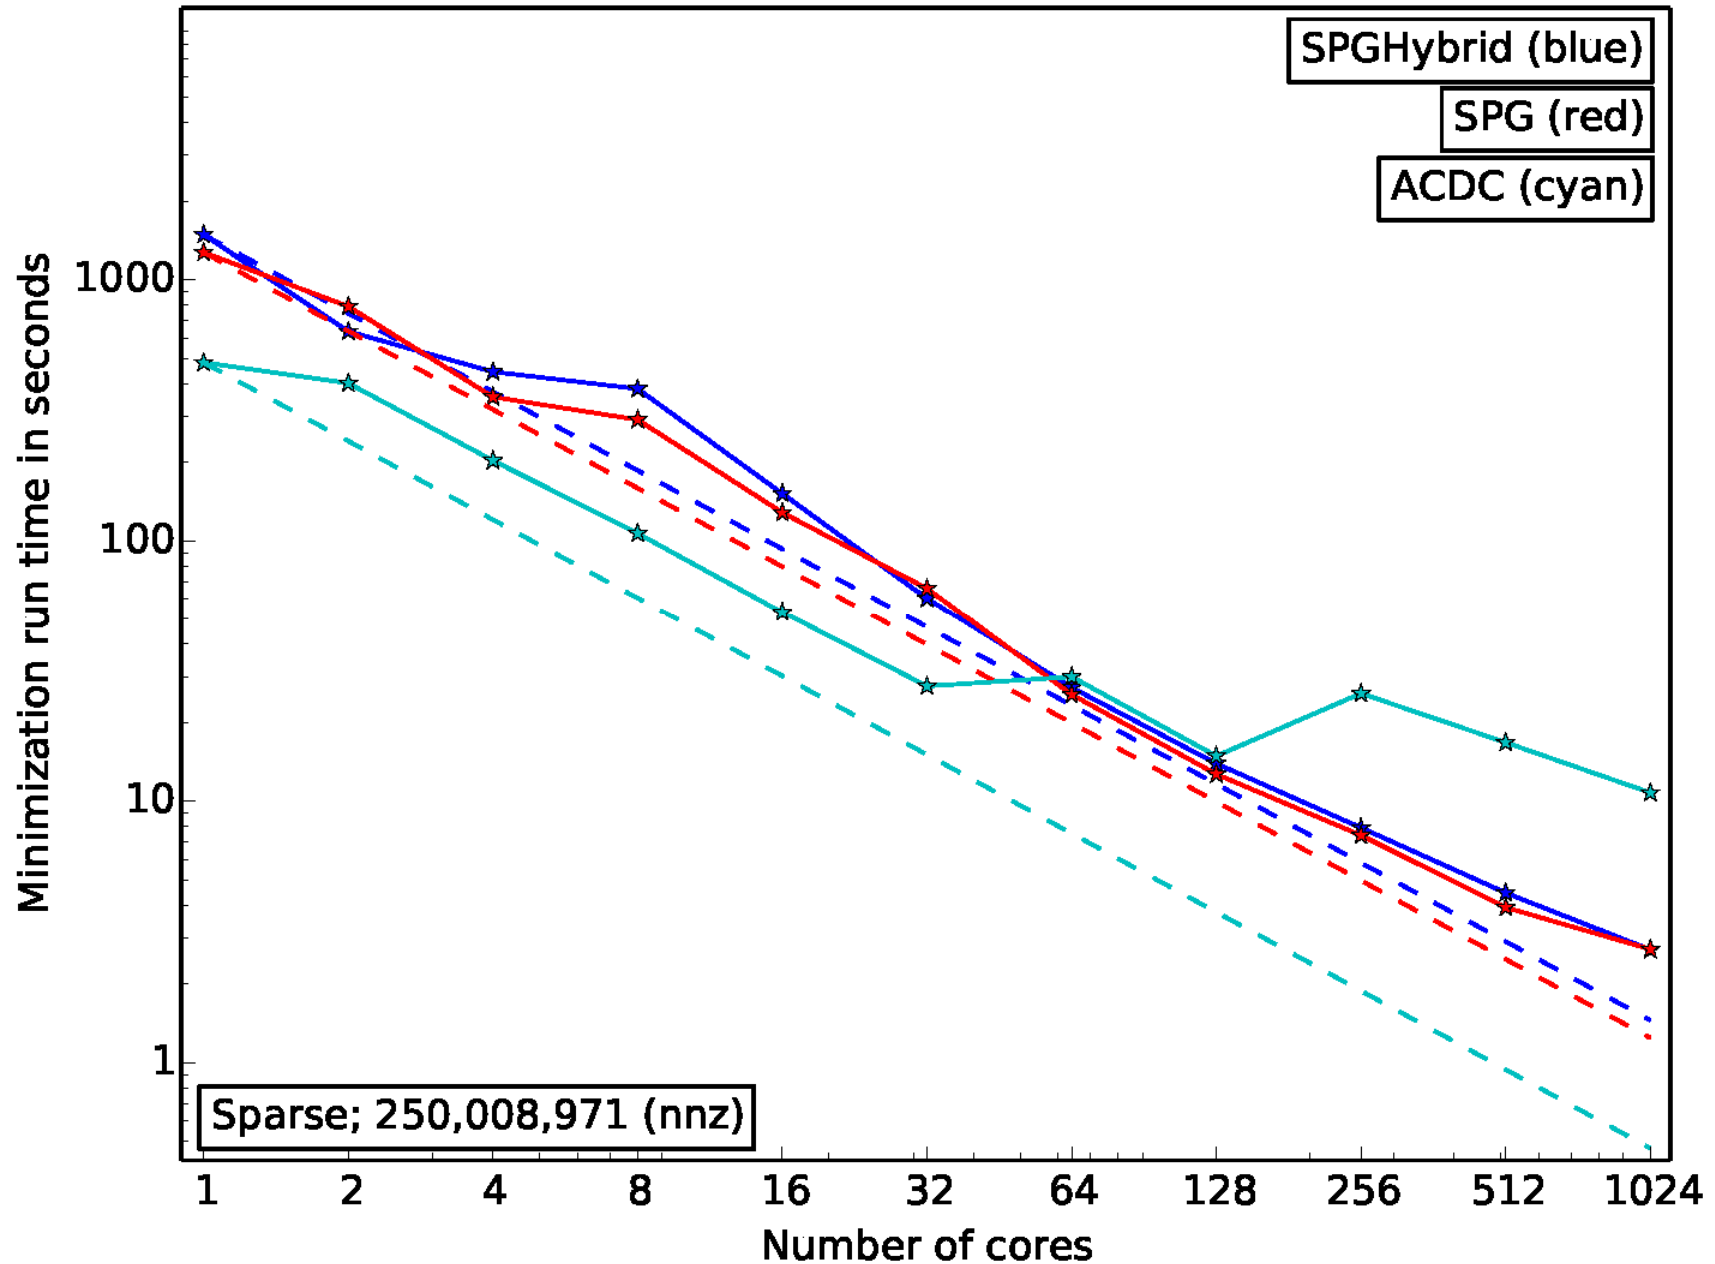
Sparse: Randomly Generated (Fairly Large)

Random Problem on Stampede (Intel Xeon e5-2680 @ 2.7 GHz)



Sparse: Randomly Generated (Varying Parameters for ACDC)

Random Problem on Stampede (Intel Xeon e5-2680 @ 2.7 GHz)



Discussion of Experimental Results

- Dense code scales well
- Balanced nonzero distribution seems to work best and scales quite well
- But only beats heuristic column balancing by a small amount (will this remain true for larger real datasets?)
 - However, it also has lower setup time (not shown here), so it's preferable to heuristic column balancing
- Epetra does not scale well on these datasets 😞
- Not dominated by ACDC in a parallel setting
 - ACDC specifically designed for Lasso-like problems
 - Our fundamental method is generic...
 - ... yet reasonably competitive in parallel

The “Object-Parallel” Approach...

... allowed application development to focus on efficiency of objective/gradient evaluations

- But not more complicated things like structure of matrix factors
- We could organize the data to optimize efficiency of the most compute-intensive operations
- We tried four different data representations and two different data layouts for x and y ...
- ... but the SPG code remained identical
- Parallelism applied to underlying operations of serial method: iteration count largely independent of number of processors.

Same principles could apply to problems with general constraints, through approximate augmented Lagrangian method

- But also need to consider application constraint structure

The “Object-Parallel” Approach: Summary

Approach to parallel optimization by combining several well-established principles:

- Efficient first-order methods (continuing to develop)
- Object-oriented programming
- And eventually, augmented Lagrangians (for problems with general constraints)
 - ... with new “loose” approximation criteria