

# On 5<sup>th</sup> Generation of Local Training Methods in Federated Learning

Peter Richtárik



**Scientific Computing and Machine Learning Workshop (SCML)**

KAUST

November 14-18, 2022



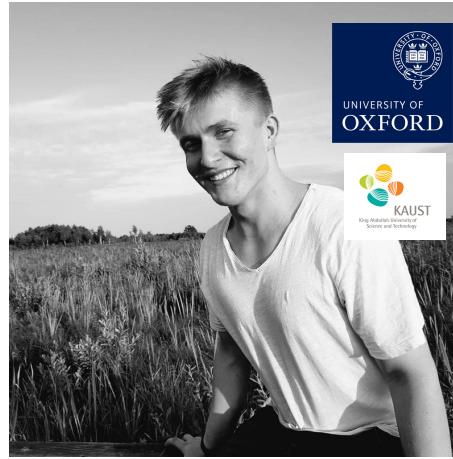
**Konstantin Mishchenko**



**Arto Maranjyan**



**Dmitry Kovalev**



**Michal Grudzien**



**Laurent Condat**



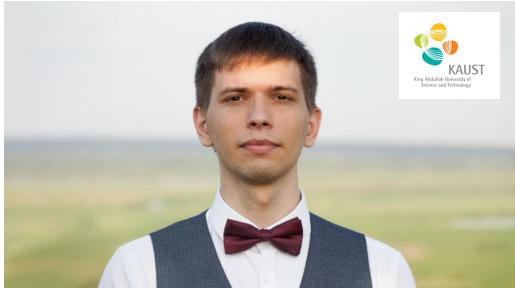
**Sebastian Stich**



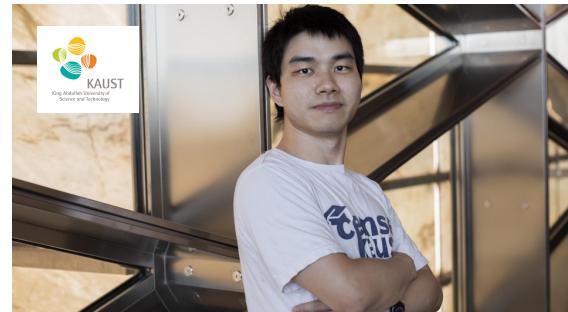
**Ivan Agarský**



**Mher Safaryan**



**Grigory Malinovsky**



**Kai Yi**



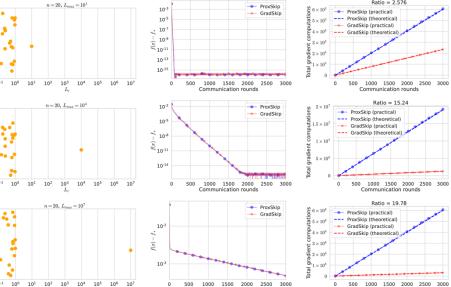
**Abdurakhmon Sadiev**

# Coauthors



# Outline of the Talk

1. Local Training
2. Brief History of Local Training
3. 5th Generation of Local Training Methods
4. ProxSkip
5. GradSkip



**Algorithm 1** ProxSkip

```

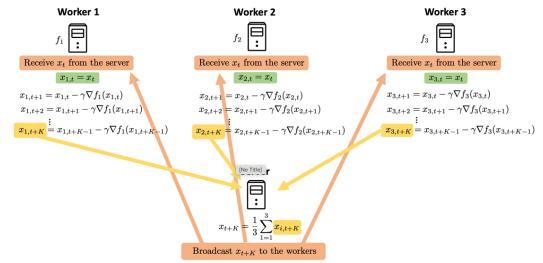
1: stepsize  $\gamma > 0$ , probability  $p > 0$ , initial iterate  $x_0 \in \mathbb{R}^d$ , initial control variate  $h_0 \in \mathbb{R}^d$ , number of iterations  $T \geq 1$ 
2: for  $t = 0, 1, \dots, T - 1$  do
3:    $\hat{x}_{t+1} = x_t - \gamma(\nabla f(x_t) - h_t)$ 
4:   Flip a coin  $\theta_t \in \{0, 1\}$  where  $\text{Prob}(\theta_t = 1) = p$ 
5:   if  $\theta_t = 1$  then
6:      $x_{t+1} = \text{prox}_{\frac{1}{p}\psi}(\hat{x}_{t+1} - \frac{\gamma}{p}h_t)$ 
7:   else
8:      $x_{t+1} = \hat{x}_{t+1}$ 
9:   end if
10:   $h_{t+1} = h_t + \frac{p}{\gamma}(x_{t+1} - \hat{x}_{t+1})$ 
11: end for

```

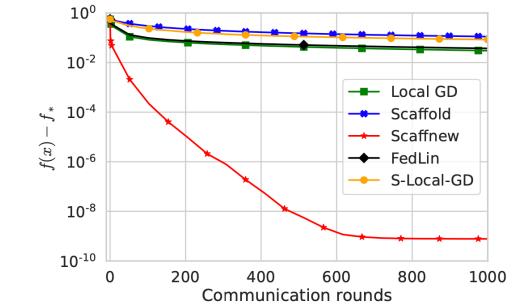
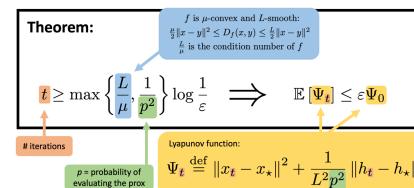


## Distributed Local Gradient Descent

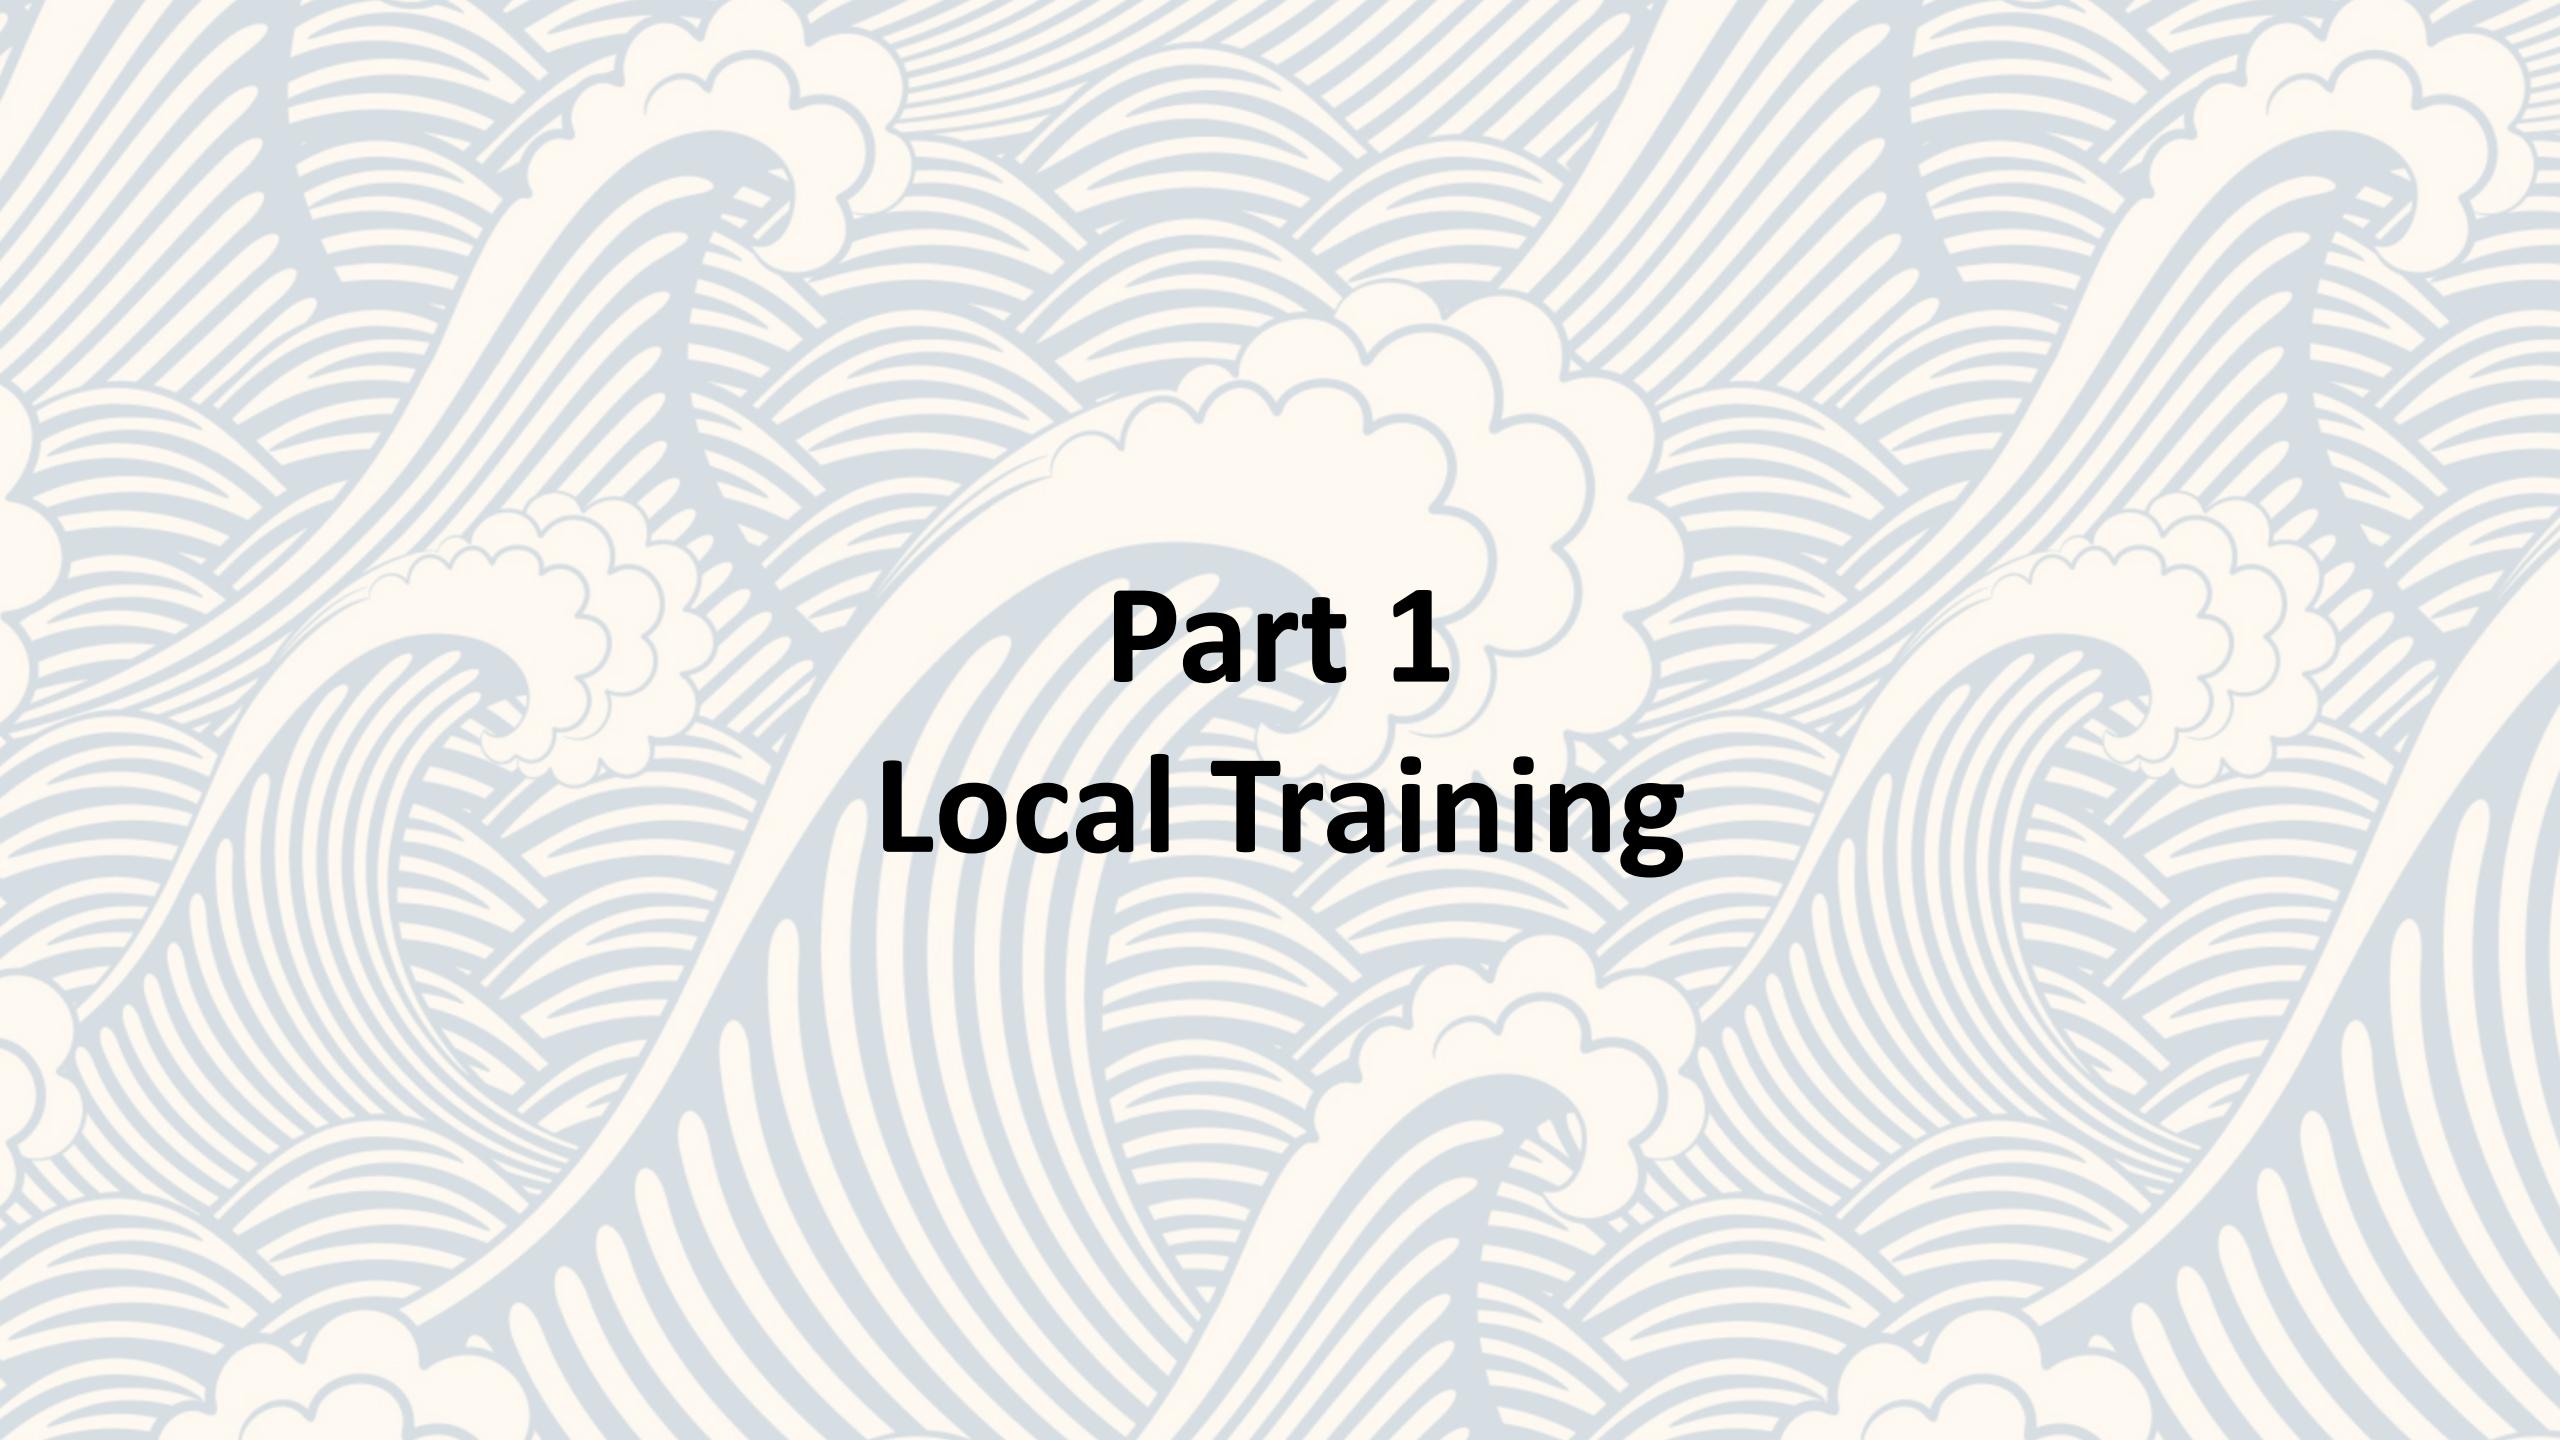
(Each worker performs  $K$  GD steps using its local function, and the results are averaged)



## ProxSkip: Bounding the # of Iterations



(c) theoretical hyper-parameters



# **Part 1**

# **Local Training**

# Optimization Formulation of Federated Learning

$$\min_{x \in \mathbb{R}^d} f(x) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n f_i(x)$$

# model parameters / features

# devices / machines

Loss on local data  $\mathcal{D}_i$  stored on device  $i$

$$f_i(x) = \mathbb{E}_{\xi \sim \mathcal{D}_i} f_{i,\xi}(x)$$

The datasets  $\mathcal{D}_1, \dots, \mathcal{D}_n$  can be arbitrarily heterogeneous

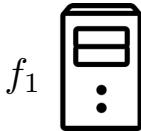
Optimization problem:

$$\min_{x \in \mathbb{R}^d} f(x) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n f_i(x)$$

# Distributed Gradient Descent

(Each worker performs 1 GD step using its local function, and the results are averaged)

Worker 1

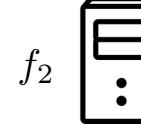


Receive  $x_t$  from the server

$$x_{1,t} = x_t$$

$$x_{1,t+1} = x_{1,t} - \gamma \nabla f_1(x_{1,t})$$

Worker 2

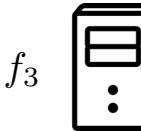


Receive  $x_t$  from the server

$$x_{2,t} = x_t$$

$$x_{2,t+1} = x_{2,t} - \gamma \nabla f_2(x_{2,t})$$

Worker 3



Receive  $x_t$  from the server

$$x_{3,t} = x_t$$

$$x_{3,t+1} = x_{3,t} - \gamma \nabla f_3(x_{3,t})$$

Server



$$x_{t+1} = \frac{1}{3} \sum_{i=1}^3 x_{i,t+1}$$

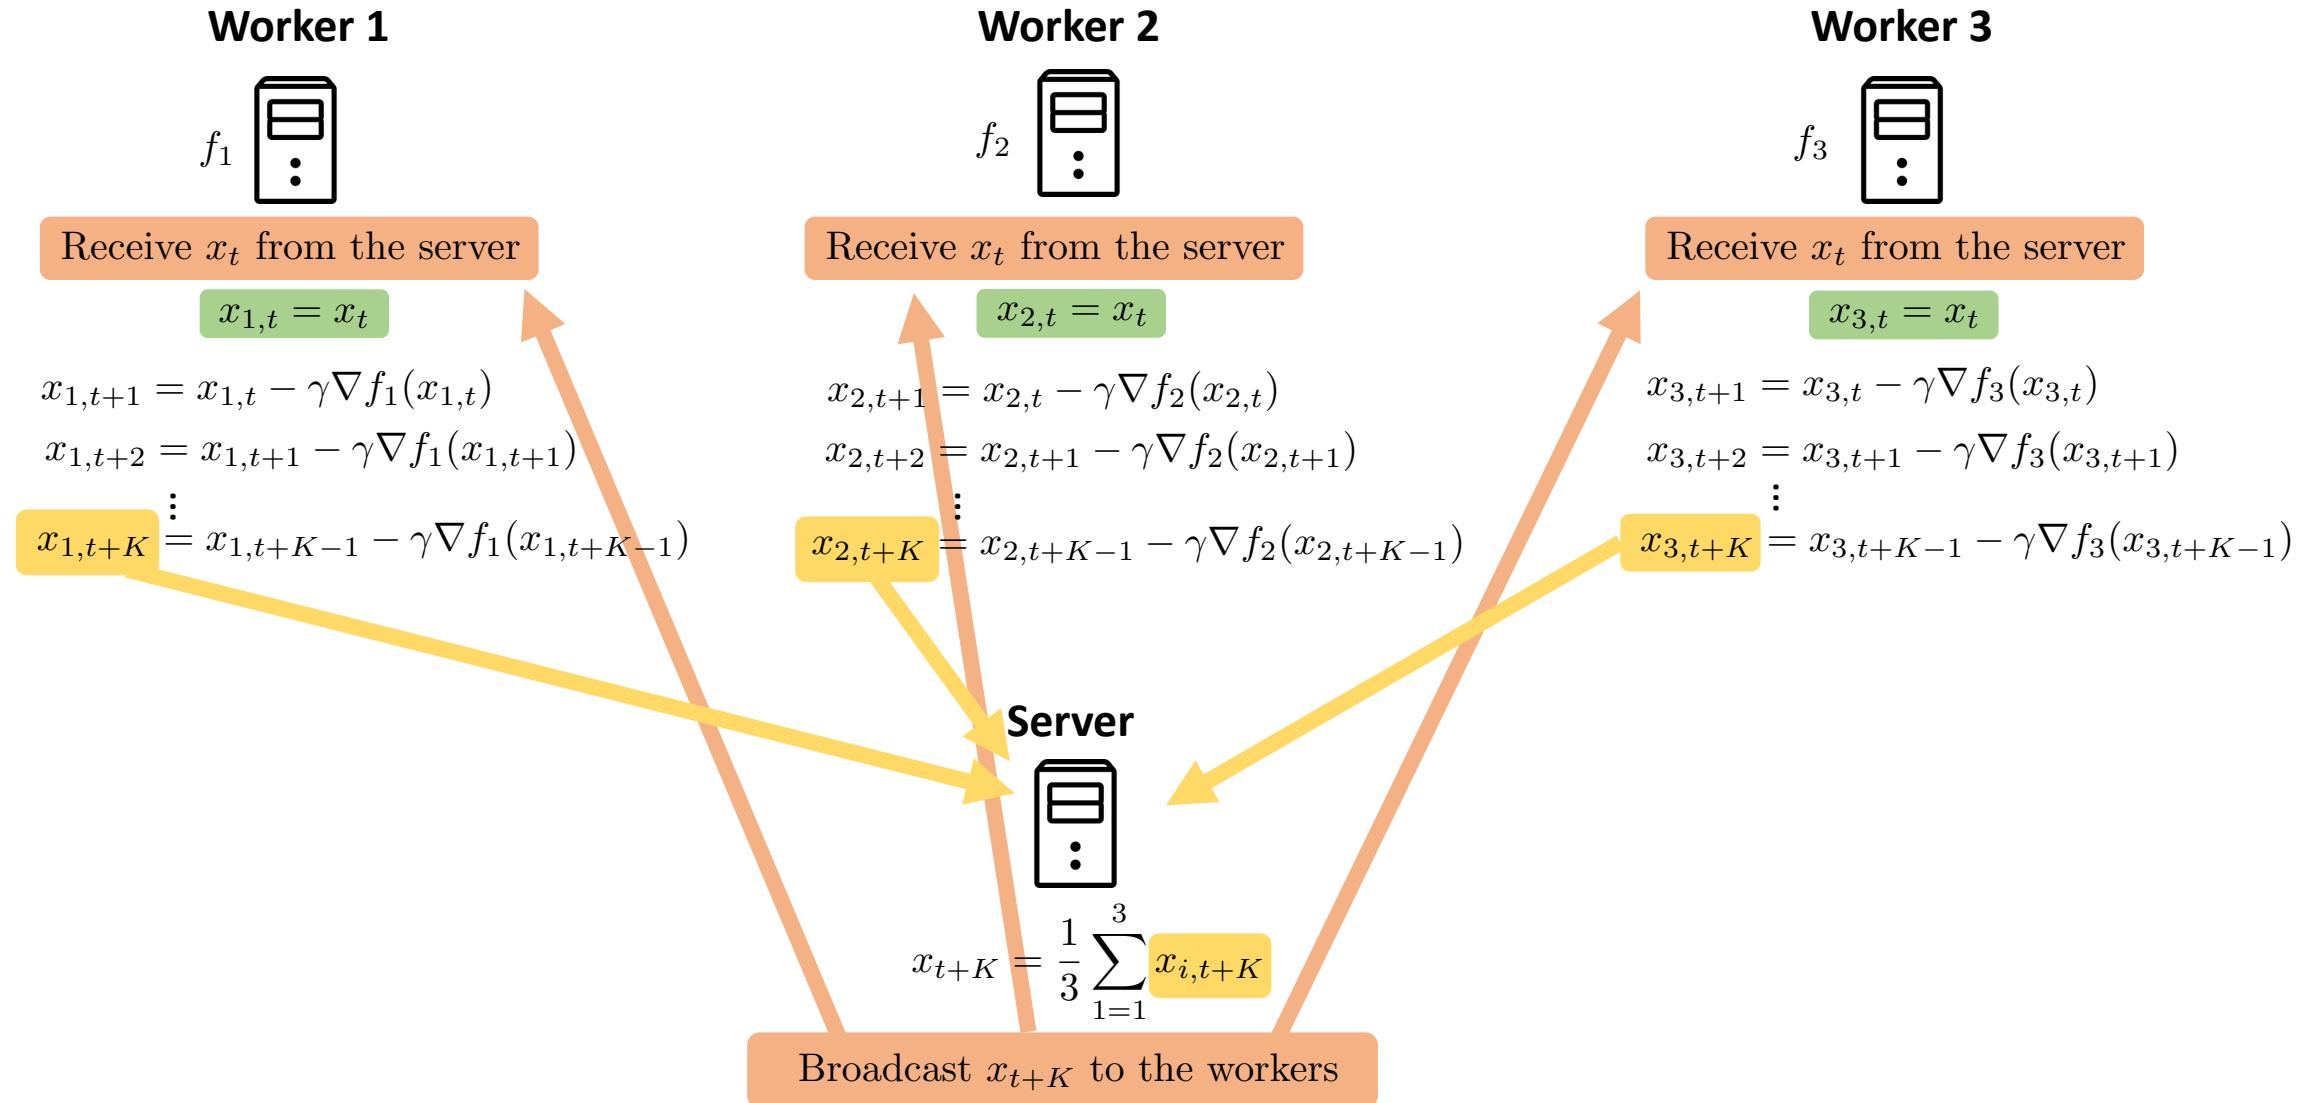
Broadcast  $x_{t+1}$  to the workers

# Distributed Local Gradient Descent

(Each worker performs  $K$  GD steps using its local function, and the results are averaged)

Optimization problem:

$$\min_{x \in \mathbb{R}^d} f(x) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n f_i(x)$$



# Part 2

# Brief History of Local Training



Grigory Malinovsky, Kai Yi and P.R.  
**Variance reduced ProxSkip: algorithm, theory and application to federated learning**  
NeurIPS 2022

# Brief History of Local Training Methods

Table 1: Five generations of local training (LT) methods summarizing the progress made by the ML/FL community over the span of 7+ years in the understanding of the *communication acceleration properties of LT*.

Generation <sup>(a)</sup>	Theory	Assumptions	Comm. Complexity <sup>(b)</sup>	Selected Key References
1. Heuristic	✗	—	empirical results only	LocalSGD [Povey et al., 2015]
	✗	—	empirical results only	SparkNet [Moritz et al., 2016]
	✗	—	empirical results only	FedAvg [McMahan et al., 2017]
2. Homogeneous	✓	bounded gradients	sublinear	FedAvg [Li et al., 2020b]
	✓	bounded grad. diversity <sup>(c)</sup>	linear but worse than GD	LFGD [Haddadpour and Mahdavi, 2019]
3. Sublinear	✓	standard <sup>(d)</sup>	sublinear	LGD [Khaled et al., 2019]
	✓	standard	sublinear	LSGD [Khaled et al., 2020]
4. Linear	✓	standard	linear but worse than GD	Scaffold [Karimireddy et al., 2020]
	✓	standard	linear but worse than GD	S-Local-GD [Gorbunov et al., 2020a]
	✓	standard	linear but worse than GD	FedLin [Mitra et al., 2021]
5. Accelerated	✓	standard	linear & better than GD	ProxSkip/Scaffnew [Mishchenko et al., 2022]
	✓	standard	linear & better than GD	ProxSkip-VR [THIS WORK]

(a) Since client sampling (CS) and data sampling (DS) can only *worsen* theoretical communication complexity, our historical breakdown of the literature into 5 generations of LT methods focuses on the full client participation (i.e., no CS) and exact local gradient (i.e., no DS) setting. While some of the referenced methods incorporate CS and DS techniques, these are irrelevant for our purposes. Indeed, from the viewpoint of communication complexity, all these algorithms enjoy best theoretical performance in the no-CS and no-DS regime.

(b) For the purposes of this table, we consider problem (1) in the *smooth* and *strongly convex* regime only. This is because the literature on LT methods struggles to understand even in this simplest (from the point of view of optimization) regime.

(c) *Bounded gradient diversity* is a uniform bound on a specific notion of gradient variance depending on client sampling probabilities. However, this assumption (as all homogeneity assumptions) is very restrictive. For example, it is not satisfied the standard class of smooth and strongly convex functions.

(d) The notorious FL challenge of handling non-i.i.d. data by LT methods was solved by Khaled et al. [2019] (from the viewpoint of *optimization*). From generation 3 onwards, there was no need to invoke any data/gradient homogeneity assumptions. Handling non-i.i.d. data remains a challenge from the point of view of *generalization*, typically by considering *personalized* FL models.



Grigory Malinovsky, Kai Yi and P.R.

Variance Reduced ProxSkip: Algorithm, Theory and Application to Federated Learning

NeurIPS 2022

# Brief History of Local Training Methods

## Generation 1: Heuristic

“No theory”

10/2014



Daniel Povey, Xiaohui Zhang, and Sanjeev Khudanpur  
**Parallel Training of DNNs with Natural Gradient and Parameter Averaging**  
*ICLR Workshops 2015*

11/2015



Philipp Moritz, Robert Nishihara, Ion Stoica, Michael I. Jordan  
**SparkNet: Training Deep Networks in Spark**  
*ICLR 2015*

02/2016



H. Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, Blaise Agüera y Arcas  
**Communication-Efficient Learning of Deep Networks from Decentralized Data**  
*AISTATS 2017*

Accepted as a workshop contribution at ICLR 2015

PARALLEL TRAINING OF DNNs WITH NATURAL GRADIENT AND PARAMETER AVERAGING

Daniel Povey, Xiaohui Zhang & Sanjeev Khudanpur  
Center for Language and Speech Processing & Related Language Technology Center of Excellence,  
The Johns Hopkins University, Baltimore, MD 21218, USA.  
(dpovey@gmail.com), (xiaohui.khudanpur@jhu.edu)

ABSTRACT

We describe the recent network training framework used in the Kaldi speech recognition toolkit, which is a grand unified framework for training DNNs with large amounts of training data. One of the key features of this framework is that it is designed to be as hardware agnostic as possible, we needed a way to use multiple machines rather than one. In order to do this, we developed a distributed parallel training framework which uses a different version of data parallelism: we have multiple SGD processes on separate machines, each of which has its own copy of the full dataset, and each machine sends updates to the individual machines. This is very effective for us for large-scale training of systems like Kaldi, where we have many machines and many datasets. We also describe a new implementation of natural gradient stochastic gradient descent (NG-SGD) that we have developed. We show that this method is significantly faster than standard SGD, and that it also handles well despite non-convexity of DNNs, so why NG-SGD is so helpful. The point of this paper is to describe the details of how we implemented this system, and how it can be used without requiring frequent data transfer (otherwise, this only adds up to about 4x to FGSM).

1. INTRODUCTION

Parallel training of neural networks generally makes use of some combination of model parallelism and data parallelism (Deng et al., 2017), and the natural approach to data parallelism involves creating work which uses a different version of data parallelism: we have multiple SGD processes on separate machines, each of which has its own copy of the full dataset, and each machine sends updates to the individual machines. This is very effective for us for large-scale training of systems like Kaldi, where we have many machines and many datasets. We also describe a new implementation of natural gradient stochastic gradient descent (NG-SGD) that we have developed. We show that this method is significantly faster than standard SGD, and that it also handles well despite non-convexity of DNNs, so why NG-SGD is so helpful. The point of this paper is to describe the details of how we implemented this system, and how it can be used without requiring frequent data transfer (otherwise, this only adds up to about 4x to FGSM).

In Section 2 we describe our problem setting, which is Deep Neural Networks (DNNs) applied to speech recognition. In Section 3 we describe the general idea behind the natural gradient and parameter averaging methods, and in Section 4 we describe the details of our implementation of the parallel training method. In Section 5 we describe the general idea behind the natural gradient and parameter averaging methods, and in Section 6 we describe the details of our implementation of the parallel training method. In Section 7 we describe the general idea behind the natural gradient and parameter averaging methods, and in Section 8 we describe the details of our implementation of the parallel training method.

The two versions of our NG-SGD method, a “simple” version and an “enhanced” one. Technical details are given in Appendices A and B respectively. Appendix C has background information on our DNN implementation.

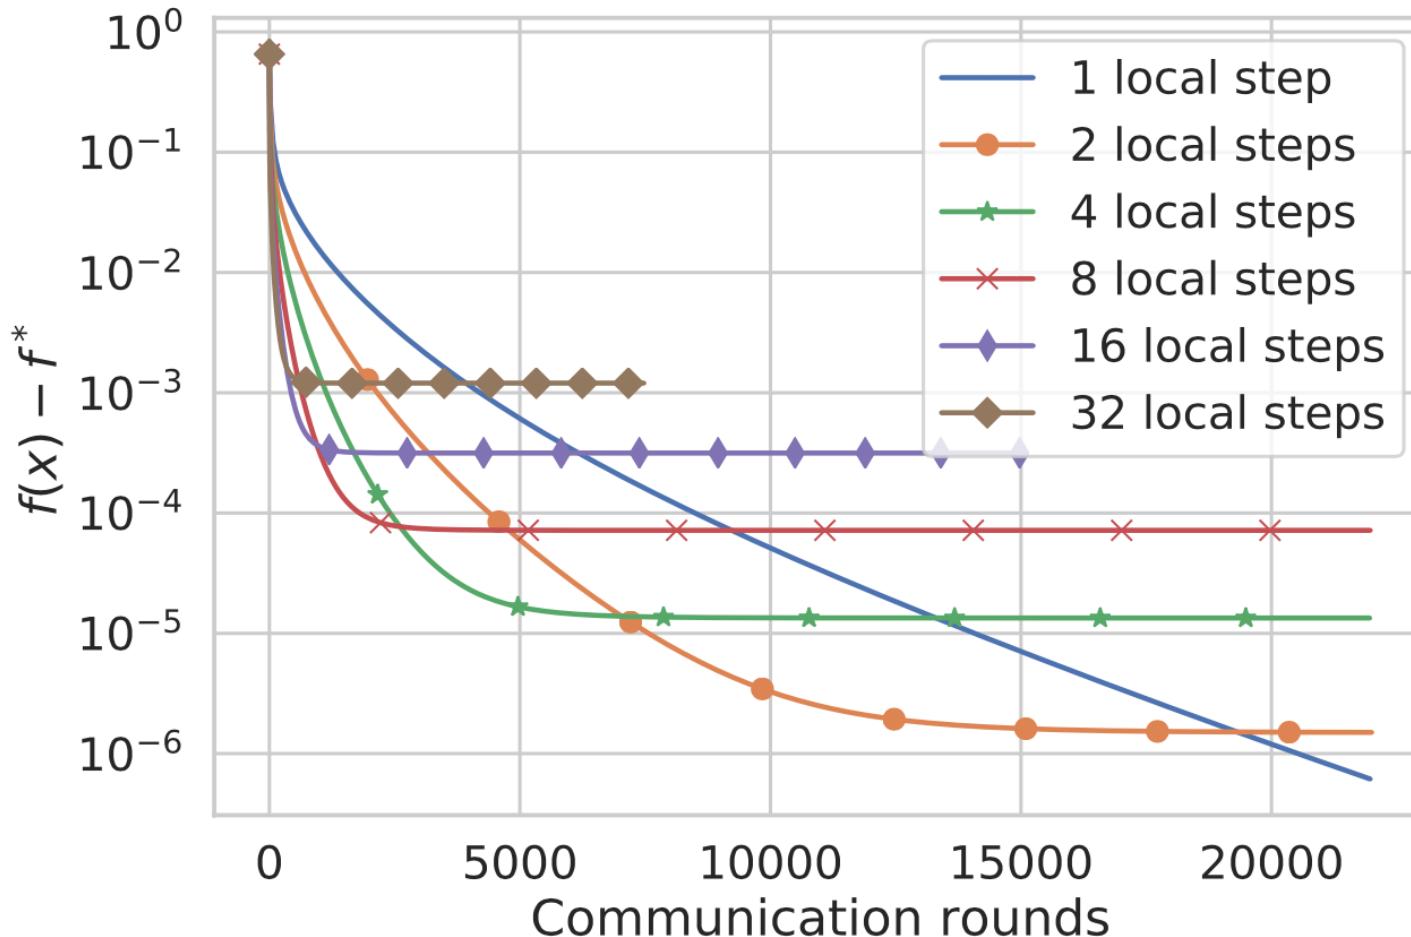
2. PROBLEM SETTING

When training DNNs for speech recognition, the immediate problem is that of classifying vectors  $x \in \mathbb{R}^n$  as corresponding to discrete labels  $y \in \mathcal{Y}$ . The dimension  $D$  is typically several hundred,

1

# Brief History of Local Training Methods

## Generation 3: Heuristic



L2-regularized logistic regression  
LibSVM mushrooms dataset

# Brief History of Local Training Methods

## Generation 2: Homogeneous

“Theory requires data to be similar/homogeneous across the clients”

07/2019



Xiang Li, Kaixuan Huang, Wenhao Yang, Shusen Wang and Zhihua Zhang  
**On the Convergence of FedAvg on Non-IID Data**  
ICLR 2020

**Bounded gradients:**

$$\|\nabla f_i(x)\| \leq B \quad \forall x \in \mathbb{R}^d \quad \forall i \in \{1, 2, \dots, n\}$$

10/2019



Farzin Haddadpour and Mehrdad Mahdavi  
**On the Convergence of Local Descent Methods in Federated Learning**  
arXiv:1910.14425, 2019

**Bounded gradient diversity (aka strong growth):**

$$\frac{1}{n} \sum_{i=1}^n \|\nabla f_i(x)\|^2 \leq C \|\nabla f(x)\|^2 \quad \forall x \in \mathbb{R}^d$$

# Brief History of Local Training Methods

## Generation 3: Sublinear

“Heterogeneous data is allowed, but the rate is worse than GD”

10/2019



Ahmed Khaled, Konstantin Mishchenko and P.R.  
**First Analysis of Local GD on Heterogeneous Data**

*NeurIPS 2019 Workshop on Federated Learning for Data Privacy and Confidentiality, 2019*

10/2019

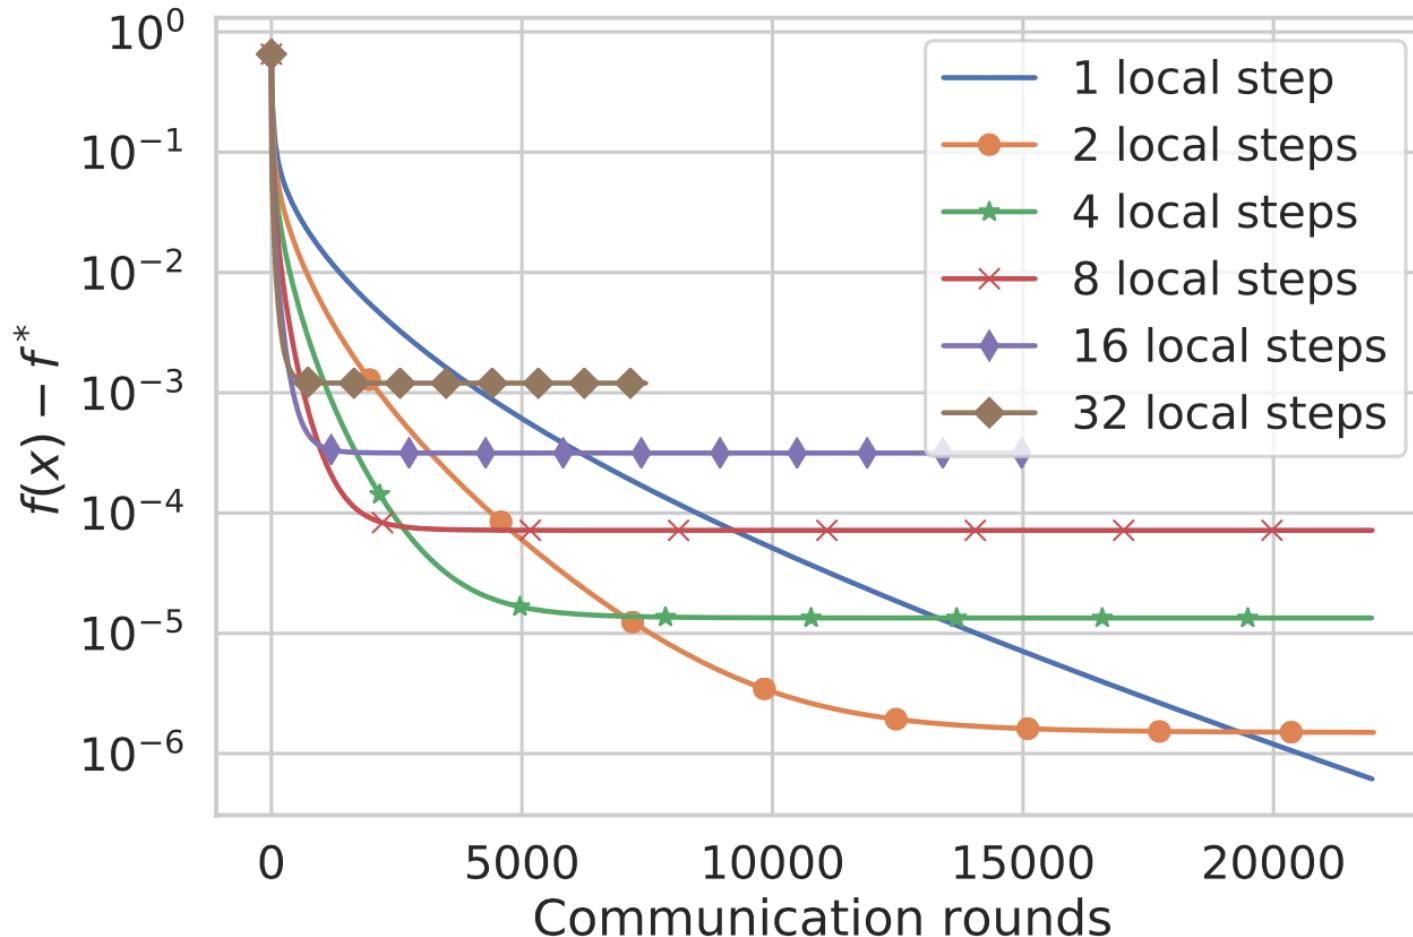


Ahmed Khaled, Konstantin Mishchenko and P.R.  
**Tighter Theory for Local SGD on Identical and Heterogeneous Data**

*AISTATS 2020*

# Brief History of Local Training Methods

## Generation 3: Sublinear



L2-regularized logistic regression  
LibSVM mushrooms dataset

# Brief History of Local Training Methods

## Generation 4: Linear

“Heterogeneous data is allowed, but the rate ay best matches that of GD”

10/2019  
Scaffold



Sai P. Karimireddy, S. Kale, M. Mohri, S. J. Reddi, S. U. Stich, A. T. Suresh  
**SCAFFOLD: Stochastic Controlled Averaging for Federated Learning**  
ICML 2020

11/2020  
S-Local-GD, Local-GD\*  
S-Local-SVRG



Eduard Gorbunov, Filip Hanzely and P.R.  
**Local SGD: Unified Theory and New Efficient Methods**  
AISTATS 2021

02/2021  
FedLin



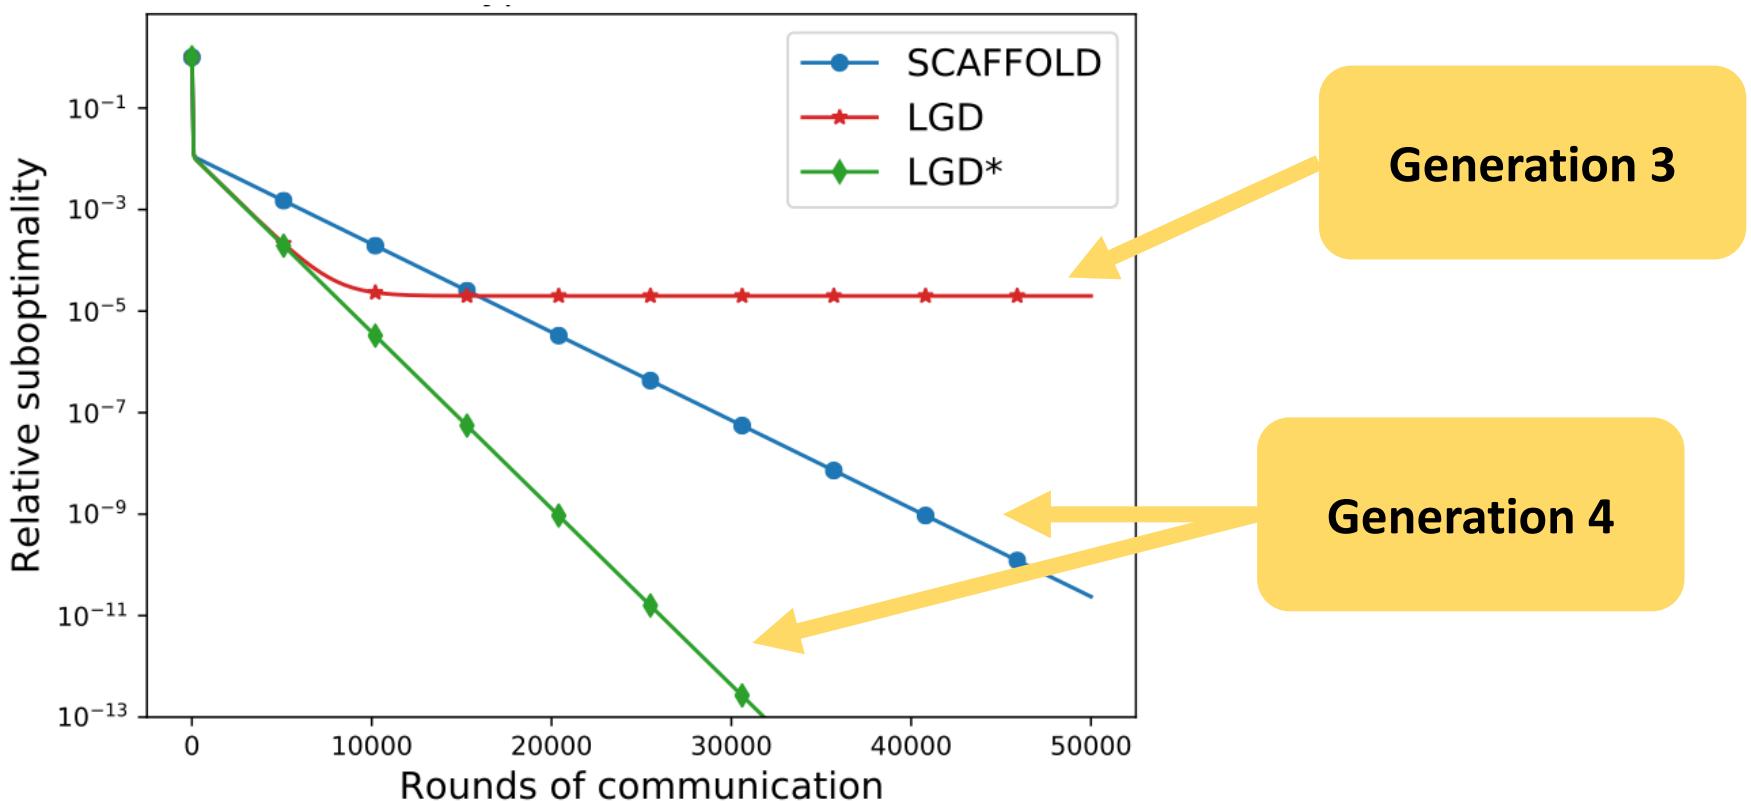
Aritra Mitra, Rayana Jaafar, George J. Pappas, Hamed Hassani  
**Linear Convergence in Federated Learning: Tackling Client Heterogeneity & Sparse Gradients**  
NeurIPS 2021

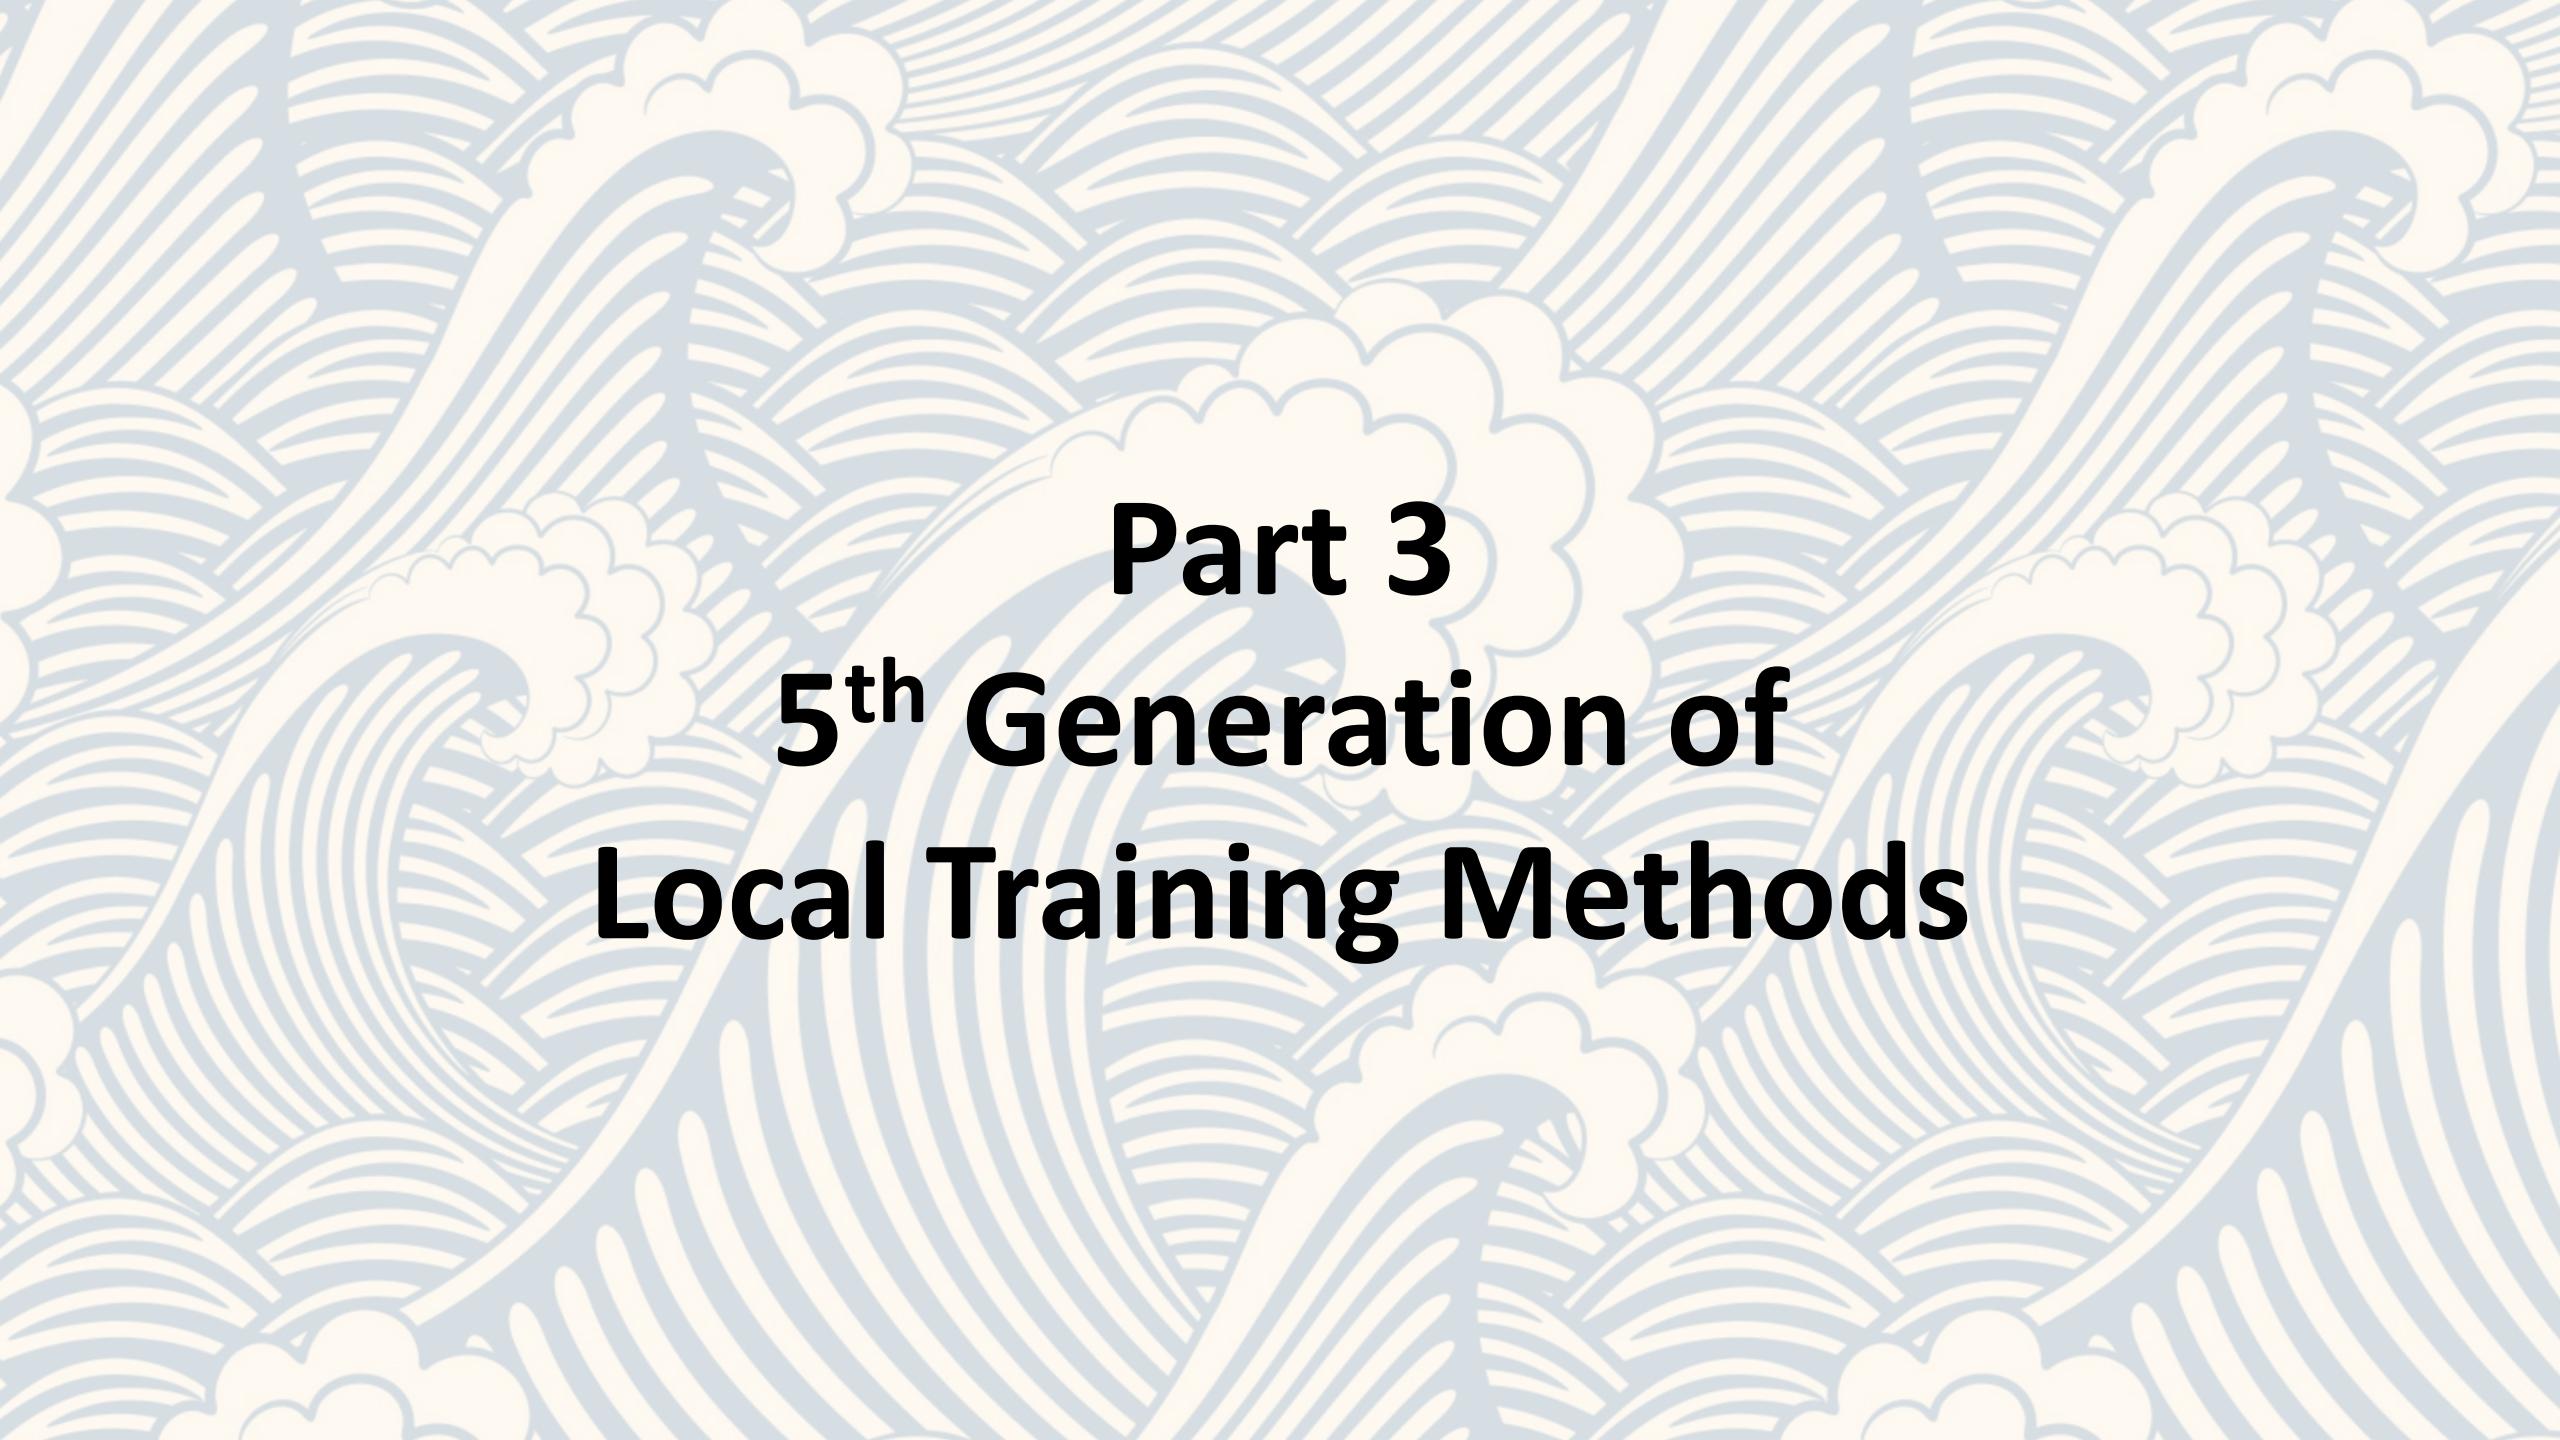
Method	$\alpha_t^*, \beta_t^*, \gamma_t^*$	Complexity	Setting	Sec
Local-SGD, Alg. 1 (Konečný et al., 2016)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$\frac{L}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}}$	UBV, Hom UBV, Het	G.1.1
Local-SGD, Alg. 1 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{(L\delta + (1-\delta)\sigma_0^2 + \gamma_0^2)\delta}{\mu^2}}$	UBV, Hom UBV, Het	G.1.1
Local-SGD, Alg. 2 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)(1-\delta)\delta}{\mu^2}}$ $+ \frac{L\delta^2\eta_{t+1}}{\mu}$	ES Hom	G.1.2
Local-SGD, Alg. 3 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)(1-\delta)\delta}{\mu^2}}$ $+ \frac{L\delta^2\eta_{t+1}}{\mu}$	ES Hom	G.1.2
Local-SGD, Alg. 4 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)(1-\delta)\delta}{\mu^2}}$ $+ \frac{L\delta^2\eta_{t+1}}{\mu}$	ES, Hom C-Het	G.1.2
Local-SGD, Alg. 5 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$m + \frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)(1-\delta)\delta}{\mu^2}}$	simple, C-Het	G.2
Local-SGD, Alg. 6 (Konečný et al., 2020)	$f_1(x_t^*) -$ $f_2(x_t^*) -$	$m + \frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)(1-\delta)\delta}{\mu^2}}$	simple, Het	G.2
S+Local-SGD, Alg. 3 (Konečný et al., 2020)	$f_1(x_t^*) - \frac{1}{m} \sum_{i=1}^m K_i^2$ $\nabla f_1(x_t^*)$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}}$	UBV, Hom UBV	G.3
S+Local-SGD, Alg. 4 (Konečný et al., 2020)	$f_1(x_t^*) - \frac{1}{m} \sum_{i=1}^m K_i^2$ $\nabla f_2(x_t^*)$	$\frac{L\delta}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}}$	UBV, Hom UBV	G.4.1
S+Local-SGD, Alg. 4 (Konečný et al., 2020)	$f_1(x_t^*) - \frac{1}{m} \sum_{i=1}^m K_i^2$ $\nabla f_2(x_t^*)$	$\frac{L}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}}$	ES, Hom	G.4.2
S+Local-SGD*, Alg. 5 (NEW)	$\nabla f_{1,i}(x_t^*) - \nabla f_{1,i}(x^*)$ $+ \nabla f_{2,i}(x_t^*) - \nabla f_{2,i}(x^*) -$	$\left( \frac{L}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}} \right) \log \frac{1}{1-\rho}$	simple, Hom	G.5
S+Local-SGD*, Alg. 6 (NEW)	$\nabla f_{1,i}(x_t^*) - \nabla f_{1,i}(x^*)$ $+ \nabla f_{2,i}(x_t^*) - \nabla f_{2,i}(x^*) -$	$\left( m + \frac{L}{\mu} + \frac{\sigma_0^2}{\mu} + \sqrt{\frac{L(\sigma_0^2 + \gamma_0^2)}{\mu}} \right) \log \frac{1}{1-\rho}$	simple, Hom	G.6

# Brief History of Local Training Methods

## Generation 4: Linear

“Heterogeneous data is allowed, but the rate ay best matches that of GD”





# **Part 3**

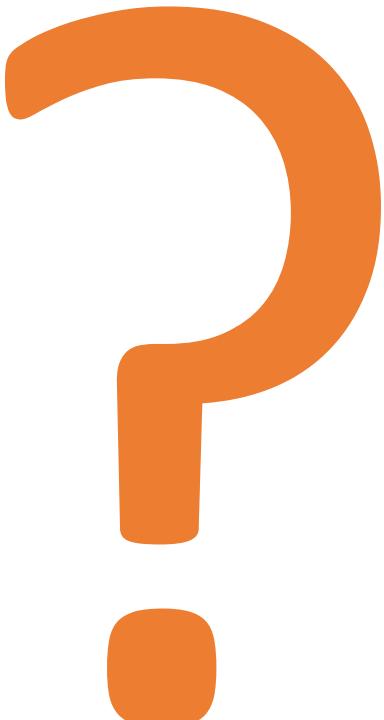
## **5<sup>th</sup> Generation of**

## **Local Training Methods**

# Brief History of Local Training Methods

## Generation 5: Accelerated

“Communication complexity is better than GD for heterogeneous data”



**In practice, local training significantly improves communication efficiency.**

**However, there is no theoretical result explaining this!**

**Is the situation hopeless, or can we show/prove that local training helps?**

# Key Property of 5<sup>th</sup> Generation Local Training Methods

Communication complexity  
of 4<sup>th</sup> generation  
local training methods

$$O\left(\frac{L}{\mu} \log \frac{1}{\varepsilon}\right)$$

Communication complexity  
of 5<sup>th</sup> generation  
local training methods

$$O\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\varepsilon}\right)$$

# ProxSkip: Yes! Local Gradient Steps Provably Lead to Communication Acceleration! Finally!<sup>†</sup>

Konstantin Mishchenko<sup>1</sup> Grigory Malinovsky<sup>2</sup> Sebastian Stich<sup>3</sup> Peter Richtárik<sup>2</sup>

## Abstract

We introduce **ProxSkip**—a surprisingly simple and provably efficient method for minimizing the sum of a smooth ( $f$ ) and an expensive nonsmooth proximable ( $\psi$ ) function. The canonical approach to solving such problems is via the proximal gradient descent (**ProxGD**) algorithm, which is based on the evaluation of the gradient of  $f$  and the prox operator of  $\psi$  in each iteration. In this work we are specifically interested in the regime in which the evaluation of prox is costly relative to the evaluation of the gradient, which is the case in many applications. **ProxSkip** allows for the expensive prox operator to be skipped in most iterations, so its iteration complexity is  $\mathcal{O}(\kappa \log^{1/\epsilon})$ , where  $\kappa$  is the condition number of  $f$ , the number of evaluations is  $\mathcal{O}(\sqrt{\kappa} \log^{1/\epsilon})$  only. Our motivation comes from federated learning, where the evaluation of the gradient operator corresponds to a local **GD** step independently on all clients, and evaluation of prox corresponds to (expensive) communication in the form of gradient exchange. In this context, **ProxSkip** offers a provable and large acceleration of communication compared to other local gradient-type methods as **FedAvg**, **SCAFFOLD**, **S-Local-GD** and others, whose theoretical communication complexity is worse than, or at best matching, that of **GD** in the heterogeneous data regime, without provable and large improvement with heterogeneity-bounding assumptions.

where  $f: \mathbb{R}^d \rightarrow \mathbb{R}$  is a smooth function, and  $\psi: \mathbb{R}^d \rightarrow \mathbb{R} \cup \{+\infty\}$  is a proper, closed and convex regularizer.

Such problems are ubiquitous, and appear in numerous applications associated with virtually all areas of science and engineering, including signal processing (Combettes & Pesquet, 2009), image processing (Luke, 2020), data science (Parikh & Boyd, 2014) and machine learning (Shalev-Shwartz & Ben-David, 2014).

### 1.1. Proximal gradient descent

<sup>†</sup> Please accept our apologies, our excitement apparently spilled over into the title. If we were to choose a more scholarly title for this work, it would be *ProxSkip: Breaking the Communication Barrier of Local Gradient Methods.*

## 1. Introduction

We study optimization problems of the form

$$\min_{x \in \mathbb{R}^d} f(x) + \psi(x), \quad (1)$$

<sup>1</sup>CNRS, ENS, Inria Sierra, Paris, France <sup>2</sup>Computer Science, King Abdullah University of Science and Technology, Thuwal, Saudi Arabia <sup>3</sup>CISPA Helmholtz Center for Information Security, Saarbrücken, Germany. Correspondence to: Peter Richtárik <peter.richtarik@kaust.edu.sa>.

<sup>†</sup> Please accept our apologies, our excitement apparently spilled over into the title. If we were to choose a more scholarly title for this work, it would be *ProxSkip: Breaking the Communication Barrier of Local Gradient Methods.*

$\text{prox}_{\gamma\psi}$ . This is the case for many regularizers, including the  $L_1$  norm ( $\psi(x) = \|x\|_1$ ), the  $L_2$  norm ( $\psi(x) = \|x\|_2^2$ ), and elastic net (Zhou & Hastie, 2005). For many further examples, we refer the reader to the books (Parikh & Boyd, 2014; Beck, 2017).

## 1.2. Expensive proximity operators

However, in this work we are interested in the situation when the evaluation of the *proximity operator is expensive*. That is, we assume that the computation of  $\text{prox}_{\gamma\psi}$  (the backward step) is costly relative to the evaluation of the gradient of  $f$  (the forward step).

A conceptually simple yet rich class of expensive proximity operators arises from regularizers  $\psi$  encoding a

# The Beginning



Konstantin Mishchenko, Grigory Malinovsky, Sebastian Stich and P.R.  
**ProxSkip: Yes! Local Gradient Steps Provably Lead to Communication Acceleration! Finally!**  
 ICML 2022

# Brief History of Local Training Methods

## Generation 5: Accelerated

“Communication complexity is better than GD for heterogeneous data”

02/2022

ProxSkip



Konstantin Mishchenko, Grigory Malinovsky, Sebastian Stich and P.R.  
**ProxSkip: Yes! Local Gradient Steps Provably Lead to Communication Acceleration! Finally!**  
*ICML 2022*

07/2022

APDA; APDA-Inexact



Abdurakhmon Sadiev, Dmitry Kovalev and P.R.  
**Communication Acceleration of Local Gradient Methods via an Accelerated Primal-Dual  
Algorithm with Inexact Prox**  
*NeurIPS 2022*

07/2022

ProxSkip-LSVRG



Grigory Malinovsky, Kai Yi and P.R.  
**Variance Reduced ProxSkip: Algorithm, Theory and Application to Federated Learning**  
*NeurIPS 2022*

07/2022

RandProx



Laurent Condat and P.R.  
**RandProx: Primal-Dual Optimization Algorithms with Randomized Proximal Updates**  
*arXiv:2207.12891, 2022*

# Brief History of Local Training Methods

## Generation 5: Accelerated

“Communication complexity is better than GD for heterogeneous data”

10/2022

GradSkip



Artavazd Maranjyan, Mher Safaryan and P.R.

**GradSkip: Communication-Accelerated Local Gradient Methods with Better Computational Complexity**

*arXiv:2210.16402, 2022*

10/2022

Compressed-Scaffnew



Laurent Condat, Ivan Agarský and P.R.

**Provably Doubly Accelerated Federated Learning: The First Theoretically Successful Combination of Local Training and Compressed Communication**

*arXiv:2210.13277, 2022*

10/2022

5GCS



Michał Grudzien, Grigory Malinovsky and P.R.

**Can 5th Generation Local Training Methods Support Client Sampling? Yes!**

*preprint, 2022*

# Brief History of Local Training Methods

## Generation 5: Accelerated

	Comm. Acceleration	Local Optimizer	# Local Training Steps	Total Complexity (Comm. + Compute)	Client Sampling?	Comm. Compression?	Supports Decentralized Setup?	Key Insight
<b>ProxSkip</b> 2/22, ICML 22	✓ $\mathcal{O}\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\epsilon}\right)$	GD	$\sqrt{\frac{L}{\mu}}$	=	✗	✗	✓	First 5th generation local training method
<b>APDA-Inexact</b> 7/22, NeurIPS 22	✓ $\mathcal{O}\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\epsilon}\right)$	any	better	better	✗	✗	✓	Can use more powerful local solvers which take fewer local GD-type steps
<b>VR-ProxSkip</b> 7/22, NeurIPS 22	✓ worse	VR-SGD	worse	better	✗	✗	✗	Running variance reduced SGD locally can lead to better total complexity than ProxSkip
<b>RandProx</b> 7/22	✓ $\mathcal{O}\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\epsilon}\right)$	GD	$\sqrt{\frac{L}{\mu}}$	=	✗	✗	✓	ProxSkip = VR mechanism for compressing the prox
<b>GradSkip</b> 10/22	✓ $\mathcal{O}\left(\sqrt{\frac{L}{\mu}} \log \frac{1}{\epsilon}\right)$	GD	better	better	✗	✗	✗	Workers containing less important data can do fewer local training steps!
<b>Compressed Scaffnew</b> 10/22	✓ worse	GD	worse	better	✗	✓	✗	Can compress uplink, leads to better overall communication complexity than ProxSkip.
<b>5GCS</b> 10/22	✓ worse	any	$\sqrt{\frac{L}{\mu}}$	worse	✓	✗	✗	Can do client sampling

# Part 4

# ProxSkip: Local Training Provably Leads to Communication Acceleration



Konstantin Mishchenko, Grigory Malinovsky, Sebastian Stich and P.R.

**ProxSkip: Yes! Local Gradient Steps Provably Lead to Communication Acceleration! Finally!**

ICML 2022

# Federated Learning: ProxSkip vs Baselines

Table 1. The performance of federated learning methods employing multiple local gradient steps in the strongly convex regime.

method	# local steps per round	# floats sent per round	stepsize on client $i$	linear rate?	# rounds	rate better than GD?
GD (Nesterov, 2004)	1	$d$	$\frac{1}{L}$	✓	$\tilde{\mathcal{O}}(\kappa)$ <sup>(c)</sup>	✗
LocalGD (Khaled et al., 2019; 2020)	$\tau$	$d$	$\frac{1}{\tau L}$	✗	$\mathcal{O}\left(\frac{G^2}{\mu n \tau \varepsilon}\right)$ <sup>(d)</sup>	✗
Scaffold (Karimireddy et al., 2020)	$\tau$	$2d$	$\frac{1}{\tau L}$ <sup>(e)</sup>	✓	$\tilde{\mathcal{O}}(\kappa)$ <sup>(c)</sup>	✗
S-Local-GD <sup>(a)</sup> (Gorbunov et al., 2021)	$\tau$	$d < \# < 2d$ <sup>(f)</sup>	$\frac{1}{\tau L}$	✓	$\tilde{\mathcal{O}}(\kappa)$	✗
FedLin <sup>(b)</sup> (Mitra et al., 2021)	$\tau_i$	$2d$	$\frac{1}{\tau_i L}$	✓	$\tilde{\mathcal{O}}(\kappa)$ <sup>(c)</sup>	✗
Scaffnew <sup>(g)</sup> (this work) for any $p \in (0, 1]$	$\frac{1}{p}$ <sup>(h)</sup>	$d$	$\frac{1}{L}$	✓	$\tilde{\mathcal{O}}\left(p\kappa + \frac{1}{p}\right)$ <sup>(c)</sup>	✓ (for $p > \frac{1}{\kappa}$ )
Scaffnew <sup>(g)</sup> (this work) for optimal $p = \frac{1}{\sqrt{\kappa}}$	$\sqrt{\kappa}$ <sup>(h)</sup>	$d$	$\frac{1}{L}$	✓	$\tilde{\mathcal{O}}(\sqrt{\kappa})$ <sup>(c)</sup>	✓

<sup>(a)</sup> This is a special case of S-Local-SVRG, which is a more general method presented in (Gorbunov et al., 2021). S-Local-GD arises as a special case when full gradient is computed on each client.

<sup>(b)</sup> FedLin is a variant with a fixed but different number of local steps for each client. Earlier method S-Local-GD has the same update but random loop length.

<sup>(c)</sup> The  $\tilde{\mathcal{O}}$  notation hides logarithmic factors.

<sup>(d)</sup>  $G$  is the level of dissimilarity from the assumption  $\frac{1}{n} \sum_{i=1}^n \|\nabla f_i(x)\|^2 \leq G^2 + 2LB^2 (f(x) - f_*)$ ,  $\forall x$ .

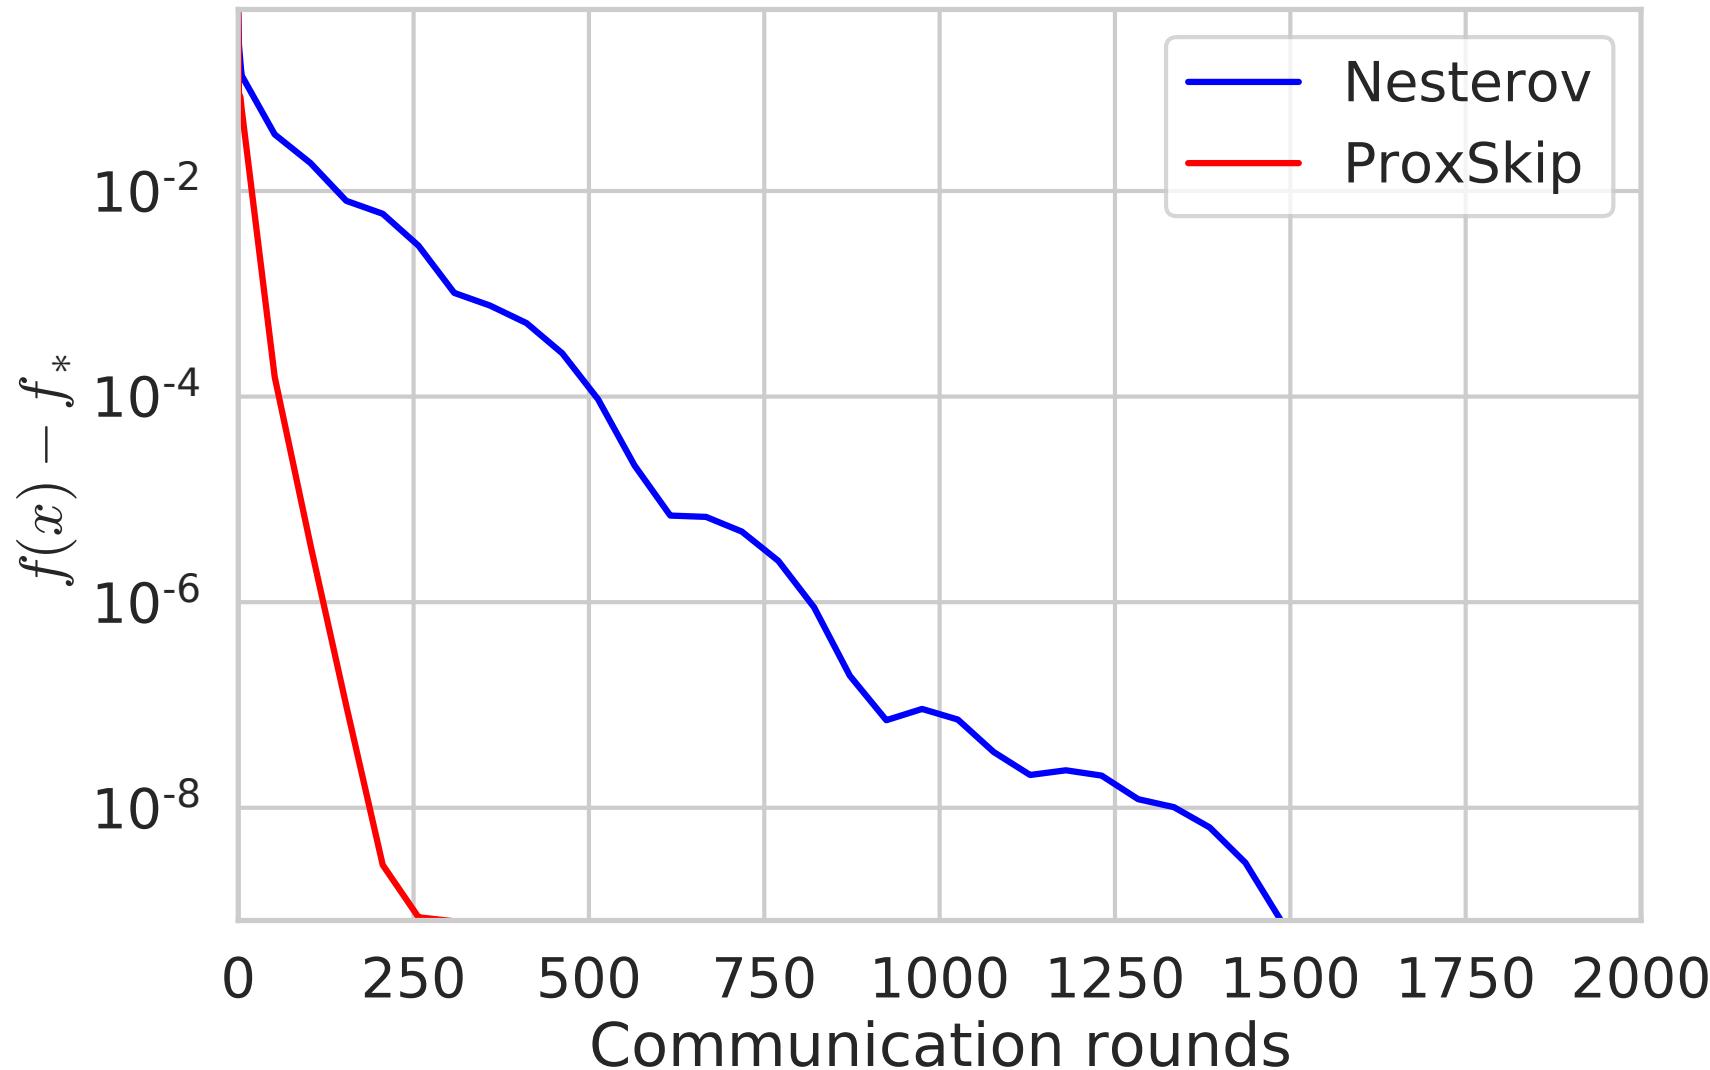
<sup>(e)</sup> We use Scaffold's cumulative local-global stepsize  $\eta_l \eta_g$  for a fair comparison.

<sup>(f)</sup> The number of sent vectors depends on hyper-parameters, and it is randomized.

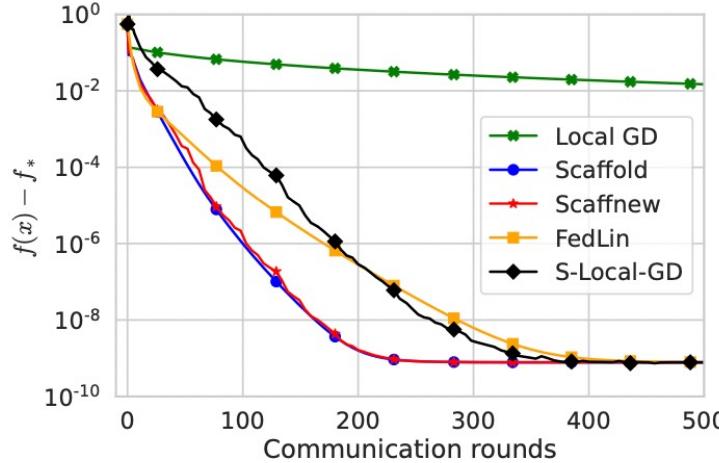
<sup>(g)</sup> Scaffnew (Algorithm 2) = ProxSkip (Algorithm 1) applied to the consensus formulation (6) + (7) of the finite-sum problem (5).

<sup>(h)</sup> ProxSkip (resp. Scaffnew) takes a *random* number of gradient (resp. local) steps before prox (resp. communication) is computed (resp. performed). What is shown in the table is the *expected* number of gradient (resp. local) steps.

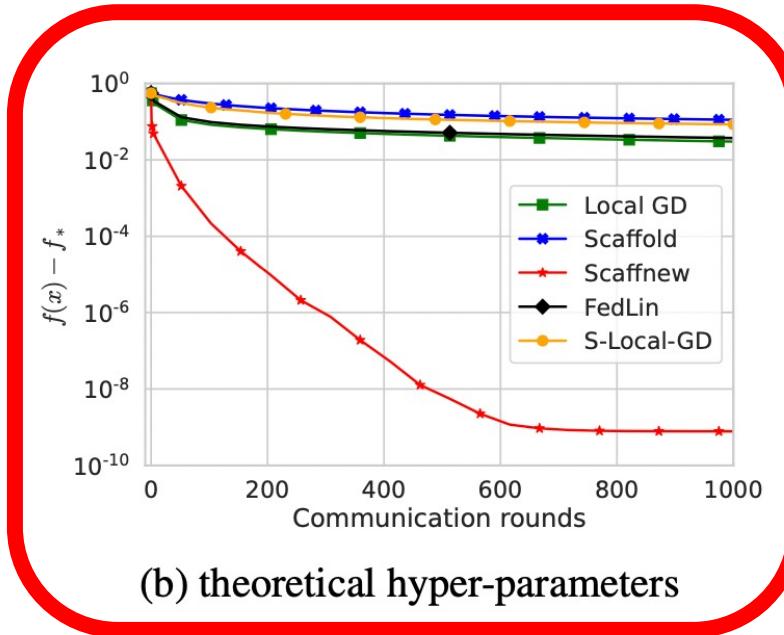
# ProxSkip vs Nesterov



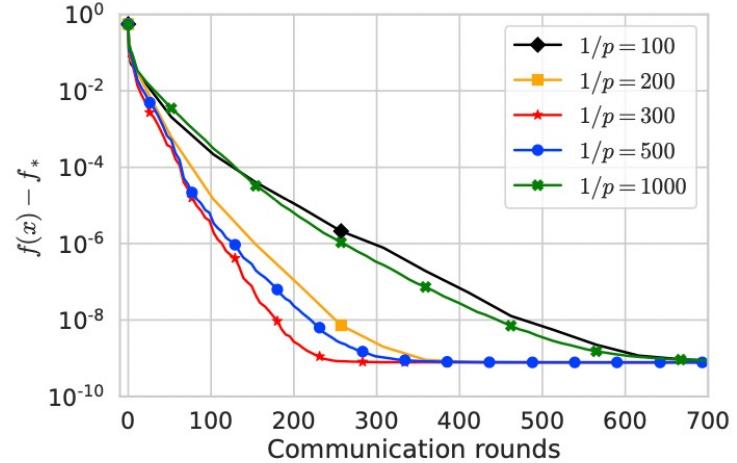
# ProxSkip + Deterministic Gradients



(a) tuned hyper-parameters



(b) theoretical hyper-parameters



(c) different options of  $p$

**Figure 1. Deterministic Case.** Comparison of [Scaffnew](#) to other local update methods that tackle data-heterogeneity and to [LocalGD](#). In (a) we compare communication rounds with optimally tuned hyper-parameters. In (b), we compare communication rounds with the algorithm parameters set to the best theoretical stepsizes used in the convergence proofs. In (c), we compare communication rounds with the algorithm stepsize set to the best theoretical stepsize and different options of parameter  $p$ .

L2-regularized logistic regression:

$$f(x) = \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-b_i a_i^\top x)) + \frac{\lambda}{2} \|x\|^2$$

$$a_i \in \mathbb{R}^d, b_i \in \{-1, +1\}, \lambda = L/10^4$$

w8a dataset from LIBSVM library (Chang & Lin, 2011)

# Consensus Reformulation

**Original problem:**  
optimization in  $\mathbb{R}^d$

$$\min_{x \in \mathbb{R}^d} \left\{ f(x) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n f_i(x) \right\}$$

**Bad:** non-differentiable

**Good:** Indicator function of a  
nonempty closed convex set

**Consensus reformulation:**  
optimization in  $\mathbb{R}^{nd}$

$$\min_{x_1, \dots, x_n \in \mathbb{R}^d} \left\{ \frac{1}{n} \sum_{i=1}^n f_i(x_i) + \psi(x_1, \dots, x_n) \right\}$$

$$\psi(x_1, \dots, x_n) \stackrel{\text{def}}{=} \begin{cases} 0, & \text{if } x_1 = \dots = x_n, \\ +\infty, & \text{otherwise.} \end{cases}$$

# Consensus Reformulation

**Original problem:**  
optimization in  $\mathbb{R}^d$

$$\min_{x \in \mathbb{R}^d} \left\{ f(x) \stackrel{\text{def}}{=} \frac{1}{n} \sum_{i=1}^n f_i(x) \right\}$$

**Bad:** non-differentiable

**Good:** proper closed convex

**Consensus reformulation:**  
optimization in  $\mathbb{R}^{nd}$

$$\min_{x_1, \dots, x_n \in \mathbb{R}^d} \left\{ \frac{1}{n} \sum_{i=1}^n f_i(x_i) + \psi(x_1, \dots, x_n) \right\}$$

$$\psi(x_1, \dots, x_n) : \mathbb{R}^{nd} \rightarrow \mathbb{R} \cup \{+\infty\}$$

is a proper closed convex function

$$\text{epi}(\psi) \stackrel{\text{def}}{=} \{(x, t) \mid \psi(x) \leq t\}$$

The epigraph of  $\psi$  is a closed and convex set

# Three Assumptions

The epigraph of  $\psi$  is a closed and convex set

$$\text{epi}(\psi) \stackrel{\text{def}}{=} \{(x, t) \in \mathbb{R}^d \times \mathbb{R} \mid \psi(x) \leq t\}$$

$$\min_{x \in \mathbb{R}^d} f(x) + \psi(x)$$

**A1**  $f$  is  $\mu$ -convex and  $L$ -smooth:

$$\frac{\mu}{2} \|x - y\|^2 \leq D_f(x, y) \leq \frac{L}{2} \|x - y\|^2$$

Bregman divergence of  $f$ :

$$D_f(x, y) \stackrel{\text{def}}{=} f(x) - f(y) - \langle \nabla f(y), x - y \rangle$$

**A2**  $\psi : \mathbb{R}^d \rightarrow \mathbb{R} \cup \{+\infty\}$  is proper, closed, and convex

**A3**

$\psi$  is proximable

The proximal operator  $\text{prox}_\psi : \mathbb{R}^d \rightarrow \mathbb{R}^d$  defined by

$$\text{prox}_\psi(x) \stackrel{\text{def}}{=} \arg \min_{u \in \mathbb{R}^d} \left( \psi(u) + \frac{1}{2} \|u - x\|^2 \right)$$

can be evaluated exactly (e.g., in closed form)

# Key Method: Proximal Gradient Descent

proximal operator:

$$\text{prox}_\psi(x) \stackrel{\text{def}}{=} \arg \min_{u \in \mathbb{R}^d} \left( \psi(u) + \frac{1}{2} \|u - x\|^2 \right)$$

$$x_t - \gamma \nabla f(x_t)$$

stepsize

gradient operator

$$x \mapsto x - \gamma \nabla f(x)$$

# Proximal Gradient Descent: Theory

Theorem:

$f$  is  $\mu$ -convex and  $L$ -smooth:  
 $\frac{\mu}{2}\|x - y\|^2 \leq D_f(x, y) \leq \frac{L}{2}\|x - y\|^2$   
 $\frac{L}{\mu}$  is the condition number of  $f$

$$t \geq \frac{L}{\mu} \log \frac{1}{\varepsilon} \quad \Rightarrow \quad \|x_t - x_\star\|^2 \leq \varepsilon \|x_0 - x_\star\|^2$$

# iterations

Error tolerance

$x_\star \stackrel{\text{def}}{=} \arg \min_{x \in \mathbb{R}^d} f(x) + \psi(x)$

(for stepsize  $\gamma = \frac{1}{L}$ )

# ProxSkip: Bird's Eye View

$$\min_{x \in \mathbb{R}^d} f(x) + \psi(x)$$

1

$$\hat{x}_{t+1} = x_t - \gamma (\nabla f(x_t) - h_t)$$

2a

with probability  $1 - p$  do

$$1 - p \approx 1$$

$$x_{t+1} = \hat{x}_{t+1}$$

$$h_{t+1} = h_t$$

2b

with probability  $p$  do

$$p \approx 0$$

evaluate  $\text{prox}_{\frac{\gamma}{p}\psi}(?)$

$$x_{t+1} = ?$$

$$h_{t+1} = ?$$

# ProxSkip: The Algorithm (Detailed View)

---

## Algorithm 1 ProxSkip

```
1: stepsize  $\gamma > 0$ , probability  $p > 0$ , initial iterate  $x_0 \in \mathbb{R}^d$ , initial control variate  $\mathbf{h}_0 \in \mathbb{R}^d$ , number of iterations  $T \geq 1$ 
2: for  $t = 0, 1, \dots, T - 1$  do
3:    $\hat{x}_{t+1} = x_t - \gamma(\nabla f(x_t) - \mathbf{h}_t)$            ◊ Take a gradient-type step adjusted via the control variate  $\mathbf{h}_t$ 
4:   Flip a coin  $\theta_t \in \{0, 1\}$  where  $\text{Prob}(\theta_t = 1) = p$       ◊ Flip a coin that decides whether to skip the prox or not
5:   if  $\theta_t = 1$  then
6:      $x_{t+1} = \text{prox}_{\frac{\gamma}{p}\psi}(\hat{x}_{t+1} - \frac{\gamma}{p}\mathbf{h}_t)$  ?          ◊ Apply prox, but only very rarely! (with small probability  $p$ )
7:   else
8:      $x_{t+1} = \hat{x}_{t+1}$                                          ◊ Skip the prox!
9:   end if
10:   $\mathbf{h}_{t+1} = \mathbf{h}_t + \frac{p}{\gamma}(x_{t+1} - \hat{x}_{t+1})$  ?        ◊ Update the control variate  $\mathbf{h}_t$ 
11: end for
```

---

# ProxSkip: Bounding the # of Iterations

**Theorem:**

$f$  is  $\mu$ -convex and  $L$ -smooth:  
 $\frac{\mu}{2}\|x - y\|^2 \leq D_f(x, y) \leq \frac{L}{2}\|x - y\|^2$   
 $\frac{L}{\mu}$  is the condition number of  $f$

$$t \geq \max \left\{ \frac{L}{\mu}, \frac{1}{p^2} \right\} \log \frac{1}{\varepsilon} \quad \Rightarrow \quad \mathbb{E} [\Psi_t] \leq \varepsilon \Psi_0$$

# iterations

$p$  = probability of evaluating the prox

Lyapunov function:

$$\Psi_t \stackrel{\text{def}}{=} \|x_t - x_\star\|^2 + \frac{1}{L^2 p^2} \|h_t - h_\star\|^2$$

# ProxSkip: Optimal Prox-Evaluation Probability

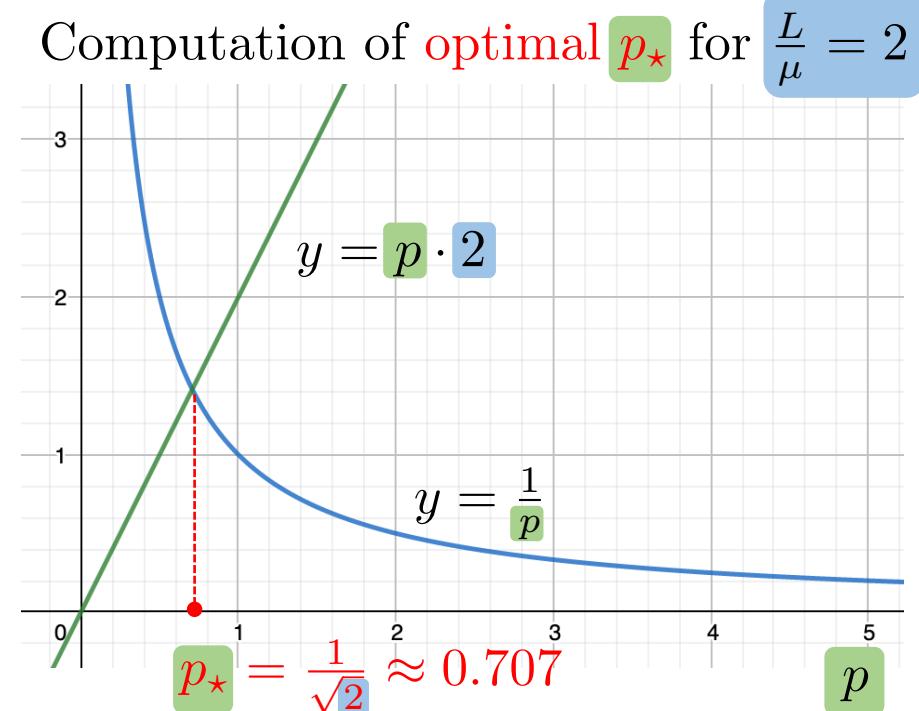
Since in each iteration we evaluate the prox with probability  $p$ ,  
the expected number of prox evaluations after  $t$  iterations is:

$$p \cdot t = p \cdot \max \left\{ \frac{L}{\mu}, \frac{1}{p^2} \right\} \cdot \log \frac{1}{\varepsilon} = \max \left\{ p \cdot \frac{L}{\mu}, \frac{1}{p} \right\} \cdot \log \frac{1}{\varepsilon}$$

$\frac{L}{\mu}$  is the condition number of  $f$

Minimized for  $p$  satisfying  $p \cdot \frac{L}{\mu} = \frac{1}{p}$

$$\Rightarrow p_{\star} = \frac{1}{\sqrt{L/\mu}}$$



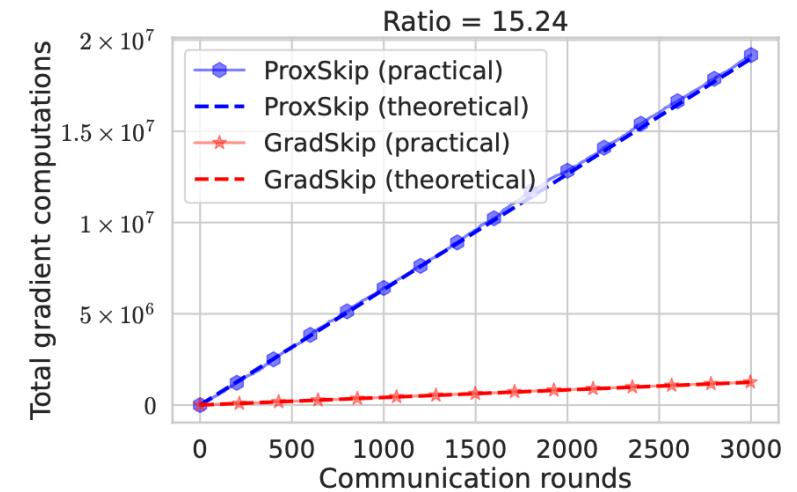
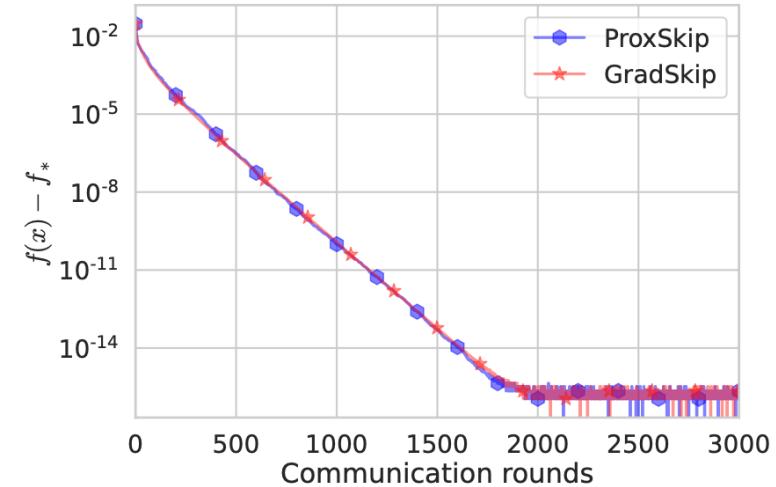
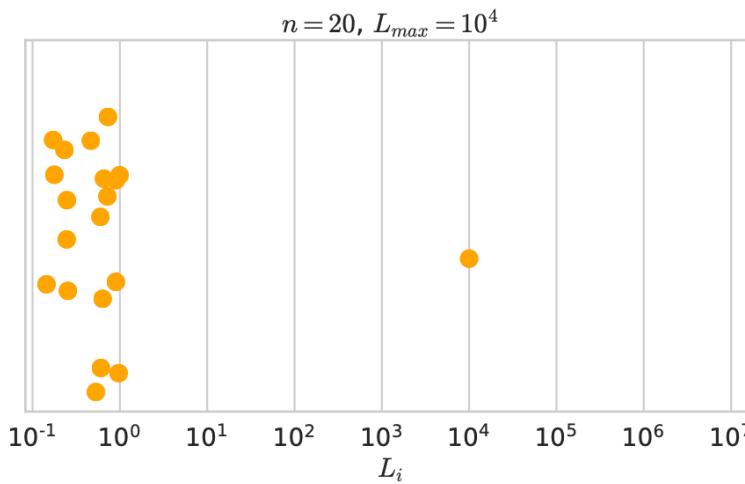
# Part 5

# GradSkip: Clients with Less Important Data can do Less Local Training



Artavazd Maranjanian, Mher Safaryan and P.R.  
**GradSkip: Communication-Accelerated Local Gradient Methods with Better  
Computational Complexity**  
*arXiv:2210.16402, 2022*

# GradSkip




---

## Algorithm 2 GradSkip+

---

```

1: Parameters: stepsize  $\gamma > 0$ , compressors  $\mathcal{C}_\omega \in \mathbb{B}^d(\omega)$  and  $\mathcal{C}_\Omega \in \mathbb{B}^d(\Omega)$ .
2: Input: initial iterate  $x_0 \in \mathbb{R}^d$ , initial control variate  $\hat{h}_0 \in \mathbb{R}^d$ , number of iterations  $T \geq 1$ .
3: for  $t = 0, 1, \dots, T-1$  do
4:    $\hat{h}_{t+1} = \nabla f(x_t) - (\mathbf{I} + \Omega)^{-1} \mathcal{C}_\Omega (\nabla f(x_t) - \hat{h}_t)$  ◊ Update the shift  $\hat{h}_{i,t}$  via shifted compression
5:    $\hat{x}_{t+1} = x_t - \gamma (\nabla f(x_t) - \hat{h}_{t+1})$  ◊ Update the iterate  $\hat{x}_{i,t}$  via shifted gradient step
6:    $\hat{g}_t = \frac{1}{\gamma(1+\omega)} \mathcal{C}_\omega (\hat{x}_{t+1} - \text{prox}_{\gamma(1+\omega)\psi} (\hat{x}_{t+1} - \gamma(1+\omega)\hat{h}_{t+1}))$  ◊ Estimate the proximal gradient
7:    $x_{t+1} = \hat{x}_{t+1} - \gamma \hat{g}_t$  ◊ Update the main iterate  $x_{i,t}$ 
8:    $\hat{h}_{t+1} = \hat{h}_{t+1} + \frac{1}{\gamma(1+\omega)} (x_{t+1} - \hat{x}_{t+1})$  ◊ Update the main shift  $h_{i,t}$ 
9: end for

```

---



Artavazd Maranjyan, Mher Safaryan and P.R.

**GradSkip: Communication-Accelerated Local Gradient Methods with Better Computational Complexity**

arXiv:2210.16402, 2022





**The End**